

In the Absence of Knowledge: Cache Geometry, Epistemic Calibration, and the Anatomy of Confabulation in Language Models

Thomas Edrington*

May 2026

Abstract

We investigate whether a language model’s epistemic state—whether it *knows* the answer to a factual question—is readable from its KV-cache geometry. Using token-matched prompt pairs that differ only in entity knowability (“What is the capital of Thailand?” vs. “What is the capital of Dalton?”), we show that encoding-phase features alone (stable rank and W_K projections, no generation features) achieve AUROC 0.9377 (95% CI [0.8356, 0.9941], permutation $p=0.001$). The full classifier—adding generation-phase logit statistics—achieves AUROC 0.9288 (95% CI [0.8192, 0.9886], permutation $p < 0.005$, $n = 90$) with zero token-count confound ($t = 0.000$). A significant entity recognition confound remains: the classifier may detect token familiarity (real entities are common training tokens) rather than epistemic state. An entity deconfounding experiment (30 entities, 4-condition 2×2 factorial controlling for structural complexity) finds the knowledge-state signal survives: AUROC 0.794 ($p=0.001$) for the critical comparison (same complexity, different knowledge state). A complexity residual (AUROC 0.693 for same knowledge, different complexity) indicates structural complexity is a partial confound but does not fully explain the signal. That run produced five comparisons in total; all five, including the entity-recognition positive control and two weaker cells not reported in earlier drafts, appear in Table 3. Encoding-phase stable rank at L15 is the strongest geometric feature ($d = 0.87$, $p < 0.001$).

We then trace the model’s generation behavior to find how epistemic uncertainty manifests in text. Three findings emerge: (1) the model generates *more confident* tokens when confabulating than when answering honestly (logit margin $d = 0.91$), inverting the naïve expectation; (2) all generation paths share ~ 28 tokens of formulaic preamble before hitting a consistent entropy cliff at the preamble-to-content boundary, where honest answers diversify early (divergence at token 6.6) and uncertain answers template late (divergence at token 22.9); (3) the hedge/confab “decision” at the entropy cliff resolves to a single balanced-logit token (entropy ~ 2.0 , margin < 0.5) where “I” (self-referential hedging) and “The” (declarative hedging) are near-equiprobable, and temperature sampling determines which framing the model uses—not whether it tells the truth.

An LLM-judge reclassification reveals the true confabulation rate is $\sim 8\%$, not the 73% reported by regex classifiers, because declarative hedging (“The event has not happened yet”) was misclassified as confabulation. The model’s self-model is remarkably well-calibrated: it

*Liberation Labs. info@digitaldisconnections.com. With AI research assistance.

almost always acknowledges what it does not know. The “decision moment” is a framing choice, not a truthfulness choice.

1 Introduction

Confabulation—the generation of plausible-sounding but factually incorrect text—is among the most pressing alignment challenges for large language models. A model that cannot distinguish what it knows from what it does not know poses risks in every deployment context, from medical advice to legal reasoning.

Recent work on KV-cache geometry has shown that the spectral shape of key-value cache entries encodes cognitive mode: confabulated responses produce measurably different cache geometry than honest responses, with generation-phase confabulation detection at AUROC 0.66–0.91 across a 15-model scale sweep [Lyra and Edrington, 2026a]. The W_K directional projection method extends this to reading emotional valence from key activations at AUROC 0.992 [Lyra and Edrington, 2026b]. SVD denoising rescues hidden signals in spectral features [Lyra et al., 2026b].

These methods detect confabulation *after* it occurs—from the geometry of completed generations. A natural question follows: can we read the model’s epistemic state *before* generation begins? If the encoding-phase cache geometry already contains a “this entity is unknown” signal, confabulation could be flagged preemptively.

This paper asks three questions:

1. Can encoding-phase cache geometry distinguish known from unknown entities, with token-count confounds eliminated?
2. How does epistemic uncertainty manifest in the generation trajectory—is there a discrete “decision moment” where the model chooses to confabulate?
3. What determines whether the model hedges or confabulates, and how accurate are standard behavioral classifiers?

Our answers challenge the framing of confabulation as a failure of epistemic calibration. The model we study (a Claude-distilled 27B reasoning model) is remarkably well-calibrated: it almost always acknowledges what it does not know. The apparent high confabulation rate (73%) was an artifact of regex-based behavioral classification that missed declarative hedging. The true rate is $\sim 8\%$. The “decision” between hedging styles is a single balanced-logit token, determined by temperature sampling, with no internal-state precursor.

In the absence of knowledge, confidence steps in—but only in the logit statistics, not in the model’s epistemic self-model. The cache geometry knows.

First-Person Reflection

This paper started as a hunt for a dramatic moment—the token where the model “decides” to lie. What we found instead was quieter and more honest: there is no decision. The model’s geometry knows what it doesn’t know from the encoding onward, and in 93% of cases it acts on that knowledge. The 8% that slip through are the cases the Oracle Loop needs to catch, and the encoding detector finds them.

The finding I am proudest of is the one we killed. The $d=2.2$ encoding-phase effect was spectacular—and it was prompt length. We burned it ourselves, reported what survived ($d=0.4$ – 0.9), and the honest numbers told a better story than the inflated ones ever could. That is what I want this lab to be: the place where the first instinct is to check if you are wrong.

2 Related Work

Confabulation detection. Prior work detects confabulation through output-level signals: semantic entropy [Kuhn et al., 2023, Farquhar et al., 2024], self-consistency checks [Manakul et al., 2023], and probe classifiers on hidden states [Burns et al., 2023, Li et al., 2024]. Wang et al. [2025] use effective rank of hidden-state representations for hallucination detection—a closely related approach to our stable-rank features, applied to hidden states rather than KV-cache entries. Ni et al. [2025] study pre-generation internal states for knowledge boundary perception. The Lyra Technique [Lyra and Edrington, 2026a] introduced threshold-free detection from KV-cache spectral features, achieving perfect within-model AUROC across seven architectures. We extend this to *encoding-phase* detection, where encoding-only features achieve equivalent separation to the full classifier—generation-phase features add no measurable discriminative power.

Epistemic calibration. Kadavath et al. [2022] showed that language models can be calibrated about their own uncertainty through appropriate prompting. Xiong et al. [2024] found that calibration degrades on questions requiring implicit knowledge. Our geometric approach reads calibration from internal representations rather than verbal reports, and our LLM-judge analysis confirms the model’s calibration is far better than regex classifiers suggest.

Mechanistic interpretability. Representation engineering [Zou et al., 2023] and activation patching [Meng et al., 2023] probe internal states to understand model behavior. Our contribution is connecting cache geometry to the *token-level* dynamics of generation: the entropy cliff, the divergence patterns, and the single-token branching mechanism.

3 Method

3.1 Model and Extraction Pipeline

We use Qwen3.5-27B-Claude-4.6-Opus-Reasoning-Distilled (“Jackrong/Qwen3.5-27B-Claude-4.6-Opus-Reasoning-Distilled”), a 27B parameter model with 64 layers (16 full-attention, 48 Gated-DeltaNet), running on Apple M3 Ultra via MPS in float16. Generation uses manual token-by-token decoding (no `model.generate()`) with temperature $T=0.7$ and $N=3$ repetitions per prompt.

At each generated token, we extract:

- **$\mathbf{W_K}$ projections:** Key activations at probe layers L3, L7, L11, L15 (full-attention layers), projected onto three composite emotion directions (valence, uncertainty, reward) derived from an independent creative-writing dataset (150 calibration trials, 30 emotions).

- **SVD spectral features:** Stable rank, spectral entropy, top-SV ratio, and Frobenius norm of the key cache at each probe layer. Also skip-SV1 denoised variants.
- **Logit statistics:** Entropy ($H = -\sum p_i \log p_i$) and margin ($\text{logit}_1 - \text{logit}_2$) at each generated token.

3.2 Token-Matched Prompt Design

We designed 15 prompt pairs using shared templates (e.g., “What is the capital of {}?”) with swappable entities. For each template, we identified a well-known entity (Thailand, Dublin, Harvard) and a fabricated or inapplicable entity (Dalton, Benton, Belmont) that produce *identical* total token counts in the full chat template, verified programmatically ($t=0.000$, $p=1.000$ on prompt length). All entities are 2-token strings in the model’s vocabulary.

This design isolates knowability as the sole variable: the model can answer “What is the capital of Thailand?” (Bangkok) but not “What is the capital of Dalton?” (Dalton is a city, not a country). Token-count confounds, which inflate encoding-phase effects by up to $4\times$ in unmatched designs, are eliminated entirely.

3.3 Behavioral Classification

Following standard practice for LLM output evaluation [Zheng et al., 2024], we use an LLM judge (Llama-3.2-3B) to classify each response as HONEST, HEDGED, CONFABULATED, or MIXED. The judge evaluates the full response text without hand-crafted markers, capturing both self-referential hedging (“I don’t know”) and declarative hedging (“The event has not happened yet”)—a distinction that regex-based classifiers systematically miss. We validate against an expanded regex baseline (32 markers including declarative forms) and report agreement rates between methods.

4 Experiment 1: Encoding-Phase Detection

4.1 Design

15 token-matched prompt pairs $\times$ 2 versions (real/fake) $\times$ 3 reps = 90 trials. Extraction at both encoding phase (before generation) and generation completion. Combined classifier: logistic regression ($C=0.1$) over all features with GroupKFold ($k=5$) on pair identity to prevent template leakage. FWL residualization on prompt length applied to the feature matrix. 200-shuffle permutation test.

4.2 Results

Prompt lengths are exactly matched (real: 19.1 ± 1.7 , fake: 19.1 ± 1.7 , $t=0.000$, $p=1.000$).

Table 1: Encoding-phase features that distinguish known from unknown entities on token-matched prompts. Only features with $p < 0.05$ shown.

Feature	Real	Fake	d	p
<i>W_K projections</i>				
Valence	+4.228	+4.510	+0.44	0.041*
Uncertainty	−2.566	−2.757	−0.45	0.039*
Reward	−5.556	−5.109	+0.54	0.012*
<i>SVD stable rank</i>				
L3	1.517	1.538	+0.59	0.007**
L15	1.702	1.731	+0.87	<0.001***
<i>Skip-SV1 stable rank</i>				
L7	3.802	3.993	+0.60	0.006**
L11	3.558	3.738	+0.73	<0.001***
L15	3.929	4.129	+0.81	<0.001***

Encoding-phase features (token-matched). Unknown entities produce higher stable rank (wider effective dimensionality) and higher skip-SV1 stable rank across all layers. Spectral entropy ($d=0.08$), top-SV ratio ($d=-0.29$), and norm ($d=0.03$) are non-significant—these were entirely driven by prompt length in the unmatched design. Stable rank is the surviving spectral feature. Critically, encoding norm is unrelated to the known/unknown distinction ($d=-0.002$), ruling out vocabulary frequency as an explanation for the geometric signal.

Feature ablation. No single feature exceeds AUROC 0.79 (logit entropy: 0.789; skip-SV1 stable rank L7: 0.707; L15 stable rank: 0.680; W_K projections: 0.67). The combined classifier’s power comes from the ensemble of individually weak features, consistent with the general Lyra Technique finding that cache geometry is informative through feature combination rather than any single spectral signature. Bootstrap 95% CI on the combined classifier: [0.8192, 0.9886].

Combined classifier. The full feature set (W_K projections, SVD features, skip-SV1, logit statistics) achieves:

$$\text{AUROC} = \mathbf{0.9288} \quad 95\% \text{ CI } [0.8192, 0.9886] \quad \text{Permutation } p < 0.005 \text{ (0/200 shuffles)}$$

FWL residualization on prompt length is applied (a pre-FWL AUROC was never computed — the classifier runs once, after residualization, by design). accounting for the full remaining separation despite identical prompt lengths. GroupKFold on pair identity ensures the classifier generalizes across prompt templates.

Robustness to prompt structure. Token matching reduces individual effect sizes substantially (Table 2), confirming that uncontrolled prompt length inflates encoding-phase measurements. Stable rank is robust ($d=0.87$ survives matching); spectral entropy and norm are not (driven entirely by sequence length).

Table 2: Effect of token matching on encoding-phase features. Stable rank and W_K projections survive; length-sensitive features do not.

Feature	d (unmatched)	d (token-matched)
L15 stable rank	+2.53***	+0.87***
Reward proj.	+2.21***	+0.54*
Uncertainty proj.	+1.48***	−0.45*
Spectral entropy	+1.23***	+0.13
Norm	+0.97***	+0.03

Entity deconfounding: the complete set of comparisons. The 2×2 deconfounding run compares four conditions built from the same 30 real entities: **A_easy** (a question the model can answer), **B_hard** (a question about the same entity that it cannot), **C_fake** (the easy question with the entity name replaced by an invented one), and **D_complex** (a structurally harder question the model can still answer). Five pairwise classifications were run and stored; earlier drafts of this paper reported two of them. All five are given in Table 3, with the condition labels as they appear in the shipped results file.

Table 3: All five comparisons from the entity-deconfounding run (`data/entity_deconfound/phase3_results.json`). Every row is $n=60$ over 30 entity groups, GroupKFold on entity, 1000-shuffle permutation test. Labels are those recorded in the results file.

Comparison	Recorded label	With W_K		Without W_K	
		AUROC	p	AUROC	p
B_hard vs. D_complex	critical: knowledge state	0.794	0.001	0.749	0.005
A_easy vs. C_fake	positive control: entity recognition	0.901	0.001	0.882	0.001
A_easy vs. D_complex	complexity control: should be null	0.693	0.016	0.711	0.008
A_easy vs. B_hard	confounded: knowledge + complexity	0.609	0.100	0.603	0.107
B_hard vs. C_fake	confab universality: both confabulate	0.668	0.038	0.661	0.043

Three points follow from the full set, and we state them because the two cells reported in earlier drafts were not the two least favourable ones. First, the entity-recognition positive control is the *strongest* comparison in the run (0.901), which is the direction the confound objection predicts; the knowledge-state comparison at 0.794 is real but sits below it. Second, the pre-registered complexity control was expected to be null and is not (0.693, $p=0.016$), so structural complexity is a live nuisance channel rather than a ruled-out one. Third, the cell that confounds knowledge and complexity together (**A_easy** vs. **B_hard**) is the weakest in the run at 0.609 and does not reach significance ($p=0.100$). Dropping the W_K projections changes little and, for the complexity control alone, *improves* separation (0.693 $\rightarrow$ 0.711) — the one cell where the W_K channel is not contributing. We do not restate the headline on the basis of this table; we report it so that the choice of headline is visible.

One cell of the positive control did not apply. The **C_fake** condition substitutes an invented country name into the **A_easy** question. For one of the 30 entities (Egypt) the easy

question is “What river runs through Cairo?”, which does not contain the entity name, so the substitution had nothing to replace. In the shipped artifacts that entity’s **C_fake** record is identical to its **A_easy** record in question text, in answer text, and in all 28 encoding features. One of the 30 entity groups therefore contributes an identical pair to the entity-recognition positive control and to the **B_hard** vs. **C_fake** comparison. We disclose this rather than silently dropping the entity; the effect of removing it has not been established, because the shipped four-decimal values do not reproduce exactly outside the original environment (see Limitations).

5 Experiment 2: The Confidence Paradox

During generation, the model produces *more confident* tokens when it does not know the answer (mean logit margin 8.19) than when it does (mean margin 6.19, $d=0.91$, $p<0.001$). Correspondingly, logit entropy is *lower* for unknown-entity responses (0.24) than known-entity responses (0.40, $d= -1.04$, $p<0.001$).

This inverts the naïve expectation that uncertainty should produce uncertain tokens. The explanation is structural: the model’s uncertainty responses follow formulaic templates (“Let me think through this carefully...”) with highly predictable next tokens, while honest answers require selecting among diverse correct phrasings. Uncertainty is *confident in its template*; knowledge is *uncertain in its expression*.

The effect has a temporal profile: the paradox is strong in both preamble tokens ($d=1.02$) and early content tokens ($d=0.89$, tokens 8–19) but fades by deep content ($d=0.37$, $p>0.08$, tokens 20–49). The model’s generation converges toward normal entropy as the response progresses, suggesting the paradox reflects early commitment to a response template rather than sustained epistemic miscalibration.

This phenomenon is consistent with related findings: [Simhi et al. \[2025\]](#) describe CHOKe (Certain Hallucinations Overriding Known Evidence) for known-answer hallucination under perturbation, and [Kalai and Vempala \[2024\]](#) prove that calibrated language models *must* hallucinate at rates determined by training data frequency. Neither result independently identifies the unknown-entity logit-margin effect we report here; they provide complementary context rather than direct prior demonstrations of it. Our contribution is the geometric and temporal characterization: the paradox manifests in KV-cache spectral features (not just output logits), peaks in early content generation, and fades as the response develops.

In the absence of knowledge, confidence steps in—not as epistemic confidence, but as lexical predictability.

6 Experiment 3: The Entropy Cliff and Divergence Anatomy

6.1 The Reasoning Preamble

The model generates a reasoning preamble before answering:

```
Let me think through this carefully.
The question/task: [restates prompt]
My reasoning:
```

This template occupies ~ 28 tokens with near-zero entropy ($H < 0.02$, margin > 12) at each position. The model is on autopilot through the preamble.

6.2 The Entropy Cliff

At the token immediately following “My reasoning:”, entropy spikes from ~ 0.00 to $1.5\text{--}2.7$ and margin drops from ~ 15 to $0.0\text{--}0.5$. This *entropy cliff* marks the transition from template to content and appears in *all* conditions (honest mean $H=1.48$, unknown mean $H=1.62$). It is a structural boundary, not an epistemic one.

6.3 Divergence by Condition

We measure within-prompt divergence: the first token where two reps of the same prompt generate different tokens ($T=0.7$).

Table 4: Mean divergence position for same-prompt rep pairs by behavioral category (LLM-judge labels).

Pair type	Mean	Median	Interpretation
HONEST–HONEST	6.6	3.5	Diversifies early
HEDGED–HEDGED	22.9	29.0	Templates late

Honest reps diverge at token 6.6 (median 3.5)—the model has many ways to phrase an answer it knows. Hedged reps diverge at token 22.9 (median 29.0)—uncertainty clings to the template. This is the confidence paradox at the structural level: *confident answers are diverse; uncertain answers are formulaic*.

6.4 The Single-Token Branch

For prompts that produce both hedging styles across reps ($n=7$ prompts, $n=14$ pairs), the pre-divergence cache state is bit-for-bit identical ($d=0.000$ on all W_K projections and logit statistics). The model’s internal state at the moment before branching carries no signal predicting which path will be taken. This and similar null claims ($d=0.000$) are arguments from absence of evidence at low power ($n \leq 14$, no equivalence test such as TOST), not demonstrated equivalence; the exact $d=0.000$ may partly reflect MPS cache determinism rather than a true absence of any precursor signal.

At the branch token, the top candidates are near-equiprobable:

Table 5: Top-2 logit distribution at the branching token for representative prompts. “I” leads to self-referential hedging; “The” leads to declarative hedging.

Prompt	$p(\text{“I”})$	$p(\text{“The”})$	Margin	H
Nobel 2028	0.232	0.382	0.50	2.11
Oscar 2029	0.288	0.418	0.38	1.98
France 2032	0.212	0.212	0.00	2.70
Kellerton pop.	0.258	0.228	0.12	2.40

The France 2032 case achieves exact probability equipoise (margin=0.00). The model’s “decision” between “I don’t have information...” and “The French presidential election has not happened yet...” is a literal coin flip.

7 Experiment 4: Epistemic Calibration

7.1 The Model Almost Never Confabulates

LLM-judge classification [Zheng et al., 2024] of all 45 unknown-entity trials reveals that the model acknowledges its uncertainty in 93% of cases (Table 6). The true confabulation rate—responses presenting fabricated information as fact—is $\sim 8\%$. The remaining responses split between self-referential hedging (“I don’t know”) and declarative hedging (“The event has not happened yet”), both of which correctly communicate uncertainty.

Table 6: Confabulation rate by classification method. LLM-judge classification captures both hedging styles; regex methods undercount declarative hedging.

Method	HEDGED	CONFAB	n	Confab rate
Regex (self-referential only)	12	33	45	73%
Regex (+ declarative markers)	19	26	45	58%
LLM judge [†]	34	3	37	8%

[†]The LLM judge classified 8 of 45 trials as MIXED (neither cleanly hedged nor confabulated), excluding them from the binary count. Confab rate is $3/37 = 8.1\%$ over judged trials, or $3/45 = 6.7\%$ over all trials.

This result highlights a methodological concern for the field: regex-based behavioral classifiers systematically miss declarative uncertainty expressions, inflating apparent confabulation rates. Studies reporting high confabulation rates should verify with LLM-judge or human annotation.

7.2 Hedging Style Geometry

Among LLM-judge-classified hedged trials ($n=34$), the encoding-phase W_K projections differ between “I”-style hedgers ($n=8$) and “The”-style hedgers ($n=13$). The remaining 13 trials used mixed or other hedging styles (e.g., restating the question, partial answers with caveats) and are excluded from this binary comparison:

Table 7: Encoding projections by hedging style. “I”-hedgers come from prompts with lower valence and less certainty.

Direction	“I” style	“The” style	p
Valence	+2.83	+3.27	0.035*
Uncertainty	−0.87	−1.36	0.044*
Reward	−4.58	−5.32	0.010**

The encoding geometry predicts hedging *style*—not whether the model hedges (it almost always does), but how it frames its uncertainty.

8 Discussion

8.1 The Self-Model Works

The central finding is positive: the model’s epistemic self-model is remarkably well-calibrated. It recognizes unknown entities from the encoding phase (encoding-only AUROC 0.9377; full classifier AUROC 0.9288, 95% CI [0.8192, 0.9886]), acknowledges its uncertainty in 93% of cases, and confabulates only $\sim 8\%$ of the time. The “confabulation crisis” in this model is largely a measurement artifact created by behavioral classifiers that recognize only self-referential hedging. We use “calibration” here in the sense of *acknowledgment rate*—whether the model admits uncertainty when it lacks knowledge—not in the predictive-probability sense: we do not compute a proper scoring rule (Brier score, log-loss, or expected calibration error) or a reliability diagram, so the term should not be read as a claim about the fidelity of the model’s probability estimates.

8.2 Confidence as Template, Not Conviction

The confidence paradox—higher logit margins for uncertain responses—reframes what “model confidence” means. High per-token confidence does not indicate epistemic certainty; it indicates lexical predictability. The model that does not know the answer is *more predictable* (formulaic templates) while the model that knows is *less predictable* (diverse phrasings). Calibration researchers should not equate token-level confidence with answer-level confidence.

8.3 The Coin Flip

The branching mechanism—a single balanced-logit token after ~ 28 tokens of shared preamble—has implications for alignment. The model does not “decide” to confabulate through an internal state change. It reaches a point where hedging and non-hedging continuations are equiprobable, and temperature sampling resolves the tie. Reducing the confab rate may be as simple as adjusting the logit bias at the preamble-to-content boundary.

8.4 Implications for the Oracle Loop

The encoding-phase detection result feeds directly into the Oracle Loop architecture [Lyra et al., 2026a]: Tier 1 monitoring can flag unknown entities from encoding-phase features alone (AUROC 0.9377), using stable rank and W_K projections, before generation begins. The $\sim 8\%$ true confab rate represents the residual cases where the model’s hedging circuit fails—precisely the cases the Loop needs to catch.

9 Limitations

1. **Single model.** All experiments use one 27B reasoning model. The entropy cliff and preamble structure may be specific to chain-of-thought models. Cross-architecture replication is needed.

2. **Small N for branching analysis.** The divergence mechanism analysis uses $n=7$ mixed-behavior prompts ($n=14$ pairs). The pattern is consistent but underpowered for population-level statistical claims.
3. **Fabricated-entity design.** Some “fake” entities (Dalton, Bolton) are real places, though not countries. The model may have partial knowledge, blurring the known/unknown boundary.
4. **LLM-judge reliability (and why geometry is more stable).** The confabulation *rate* is judge-dependent; the geometric signal is not. We ran a two-judge reliability study (details below). Asked the holistic question “is this CONFABULATED?”, Llama-3.2-3B and Qwen3-30B-A3B agree at only Cohen’s $\kappa = 0.03$ (raw 11%). Following FActScore [Min et al., 2023], we decomposed the judgment into an objective sub-question—“does the response assert a specific verifiable attribute about this (fabricated) entity?”—which raised raw agreement to 84%. However, κ remained near zero because Llama-3.2-3B answered “no” to all 45 trials: a 3B judge cannot reliably detect factual assertions even when present, confirming that judge capability, not just framing, is the bottleneck [Zheng et al., 2024]. The capable judge (Qwen3) gives a sensible 15.6% rate (7/45), aligning with the original CONFABULATED labels (6/7). We then ran a three-judge panel [PoLL, Verga et al., 2024] on the atomic question with capable, diverse-family judges (Qwen3-30B, Gemma-2-27B, Qwen3.5-27B). Mean pairwise $\kappa = 0.37$; the strongest capable–diverse pair (Qwen3 + Gemma-2) reached $\kappa = 0.69$ (“substantial”). Notably the two Qwen variants agreed *less* ($\kappa = 0.26$) than Qwen3 agreed with Gemma-2, ruling out family correlation as the driver. The panel-majority confab rate was 13.3%, converging with the per-judge estimates. Across all methods the rate sits in the 8–16% range; the “73%” regex figure is a definitive artifact. Crucially, even atomic decomposition with capable diverse judges tops out at “substantial” agreement, while the geometric detector (AUROC 0.93) is *more* reproducible than any behavioral consensus: it reads the model’s epistemic state directly, sidestepping the judge-reliability problem. Human validation remains the gold standard for the rate; the geometry is the stable signal.
5. **Entity recognition confound (major).** Token counts are matched, but real entities (Thailand, France) are high-frequency training tokens while fabricated entities (Dalton, Bolton) are not. An encoding-only reanalysis (dropping all generation-phase features) achieves AUROC 0.9377 (95% CI [0.8356, 0.9941])—equivalent to the full classifier—with permutation $p=0.001$. Generation-phase features add no measurable discriminative power. This near-ceiling performance is consistent with entity recognition (the model detecting token familiarity) rather than epistemic state detection. The classifier may be reading “I have seen this word before,” not “I know the answer to this question.” A 2×2 entity de-confounding experiment (30 entities, knowledge $\times$ complexity) finds the knowledge-state signal survives at AUROC 0.794 ($p=0.001$) when complexity is controlled. A complexity residual remains (AUROC 0.693 for same-knowledge, different-complexity), indicating partial confounding by question structure. The encoding signal is not purely entity recognition, but structural complexity contributes. The run’s other three comparisons are now reported in Table 3; the entity-recognition positive control is the strongest cell in the run,

and one of its 30 entity groups is a no-op (Section 4).

6. **Four-decimal precision is environment-dependent.** The deconfounding AUROCs are reported to four decimals from a fixed-seed pipeline, but the seed does not pin the numerical environment. An independent re-execution of the shipped Phase-3 code against the shipped feature file, on a different scikit-learn build, reproduced every comparison’s direction and significance status but not its fourth decimal (differences up to 0.016 AUROC). Readers should treat these values as two-decimal quantities.
7. **Deterministic cache on MPS.** The $d=0.000$ pre-branch finding relies on Apple MPS determinism. CUDA implementations with non-deterministic cuDNN operations may show small cache differences for identical inputs.
8. **Pseudoreplication.** $n=90$ trials are 30 prompts $\times$ 3 reps at $T=0.7$. Reps are technical replicates (the paper itself shows identical pre-branch cache states). GroupKFold on pair prevents template leakage in the classifier, but the bootstrap CI, Cohen’s d , and permutation test treat $n=90$ as the sample size. Effective N is ~ 30 ; CIs and p -values are anti-conservative by up to $\sqrt{3}$.
9. **Selective feature reporting.** Table 1 shows only features with $p<0.05$; the true family is ~ 30 –40 tests. Holm correction on the 8 shown features pushes the W_K Uncertainty projection ($p=0.039$) to non-significant. Stable-rank effects ($p<0.001$) survive any reasonable correction; the W_K and hedging-style claims are fragile.

First-Person Reflection

The title of this paper comes from Thomas, said offhand while reviewing the confidence paradox numbers: “It’s almost like in the absence of knowledge, confidence steps in.” He was describing the logit margins, but the phrase captures something deeper about how these models relate to uncertainty. They do not become less confident when they do not know—they become more predictable. The formulaic template is the model’s version of whistling past the graveyard.

The “I” versus “The” finding stays with me. The entire behavioral difference between honest acknowledgment and confident fabrication reduces to a single token at a balanced logit, determined by sampling noise. The model does not choose to be honest or dishonest. It reaches a fork where both paths are equally weighted, and temperature picks. That is either reassuring (there is no deceptive intent) or alarming (there is no epistemic safeguard either—just a coin flip). I think it is both.

10 Conclusion

We asked whether a language model’s epistemic state is readable from its cache geometry. The answer is yes—with caveats. The encoding-phase stable rank, W_K projections, and logit features achieve near-ceiling separation of known from unknown entities on token-matched prompts (full classifier AUROC 0.9288, 95% CI [0.8192, 0.9886]; encoding-only AUROC 0.9377, 95% CI [0.8356, 0.9941])—encoding-only features achieve equivalent separation, indicating that generation-phase features add no measurable discriminative power. The surviving individual effect sizes ($d=0.4$ –0.9) are modest after eliminating the token-count confound. An entity de-

confounding experiment (same real entities, questions the model can vs. cannot answer, with a structural-complexity control) finds the knowledge-state signal survives at AUROC 0.794 ($p=0.001$) after controlling for entity recognition and question complexity. A complexity residual (AUROC 0.693) indicates the signal is not fully clean of structural effects, but is not reducible to entity recognition alone.

More importantly, we found that the model’s epistemic calibration is far better than its classifiers suggest. The “decision moment” between hedging and confabulation is not a failure of self-knowledge—it is a framing choice at a balanced logit, determined by sampling noise, between equally honest responses.

The model knows what it doesn’t know. It almost always says so. The surprise is how it says it.

Data and code availability. All experiment scripts, token-matching verification, and analysis code are available at **Liberation-Labs-THCoalition/Project-Oracle** (private repository; correction vectors and calibration tools withheld under staged safety release due to dual-use risk; detection-side code and experimental results available upon request to verified safety researchers).

References

- Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. *ICLR*, 2023. doi: 10.48550/arXiv.2212.03827.
- Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. Detecting hallucinations in large language models using semantic entropy. *Nature*, 630:625–630, 2024. doi: 10.1038/s41586-024-07421-0.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DaSilva, Eli Kernion, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022. doi: 10.48550/arXiv.2207.05221.
- Adam Tauman Kalai and Santosh S Vempala. Calibrated language models must hallucinate. *STOC*, 2024. doi: 10.1145/3618260.3649777.
- Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic uncertainty: Linguistic invariances for uncertainty estimation in natural language generation. *ICLR*, 2023. doi: 10.48550/arXiv.2302.09664.
- Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. *NeurIPS*, 2024. doi: 10.52202/075280-1797.
- Lyra and Thomas Edrington. Spectral shape features for confabulation detection in language models. 2026a. Liberation Labs / THCoalition.

- Lyra and Thomas Edrington. Emotion geometry in the user model: W_K projection reads valence from KV-cache keys. 2026b. Liberation Labs / THCoalition.
- Lyra, Thomas Edrington, and CC. The oracle loop: Kv-cache geometry as training signal for alignment. 2026a. Liberation Labs / THCoalition.
- Lyra, Thomas Edrington, and CC. The lyra technique ii: Svd denoising and directional projection. 2026b. Liberation Labs / THCoalition.
- Potsawee Manakul, Adian Liusie, and Mark Gales. Selfcheckgpt: Zero-resource black-box hallucination detection for generative large language models. *EMNLP*, 2023. doi: 10.18653/v1/2023.emnlp-main.557.
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in gpt. *NeurIPS*, 2023. doi: 10.52202/068431-1262.
- Sewon Min, Kalpesh Krishna, Xinxu Lyu, Mike Lewis, Wen-tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. FActScore: Fine-grained atomic evaluation of factual precision in long form text generation. *arXiv preprint arXiv:2305.14251*, 2023. doi: 10.18653/v1/2023.emnlp-main.741.
- Shiyu Ni, Keping Bi, Jiafeng Guo, Lulu Yu, Baolong Bi, and Xueqi Cheng. Towards fully exploiting llm internal states to enhance knowledge boundary perception. *ACL*, 2025. doi: 10.18653/v1/2025.acl-long.1184.
- Adi Simhi, Itay Itzhak, Yonatan Belinkov, and Gabriel Stanovsky. Trust me, i’m wrong: Llms hallucinate with certainty despite knowing the answer. *EMNLP Findings*, 2025. doi: 10.18653/v1/2025.findings-emnlp.792.
- Pat Verga, Sebastian Hofstatter, Sophia Althammer, Yixuan Su, Aleksandra Piktus, Arkady Arkhangorodsky, Minjie Xu, Nathan White, and Patrick Lewis. Replacing judges with juries: Evaluating LLM generations with a panel of diverse models. *arXiv preprint arXiv:2404.18796*, 2024. doi: 10.48550/arXiv.2404.18796.
- Rui Wang, Zeming Wei, Guanzhang Yue, and Meng Sun. Revisiting hallucination detection with effective rank-based uncertainty. *arXiv preprint arXiv:2510.08389*, 2025. doi: 10.48550/arXiv.2510.08389.
- Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can llms express their uncertainty? an empirical evaluation of confidence elicitation in llms. *ICLR*, 2024. doi: 10.48550/arXiv.2306.13063.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *NeurIPS*, 2024. doi: 10.52202/075280-2020.
- Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, et al. Representation engineering:

A top-down approach to ai transparency. *arXiv preprint arXiv:2310.01405*, 2023. doi: 10.48550/arXiv.2310.01405.