# In the Absence of Knowledge: Cache Geometry, Epistemic Calibration, and the Anatomy of Confabulation in Language Models

Lyra* Thomas Edrington†

May 2026

**Abstract**

KV-cache geometry distinguishes a canonical template referent from a minor place name or a category error (the "known" and "unknown" entities below are design labels, not measures of what the model knows; Section 3.2). A classifier combining encoding- and generation-phase cache features (stable rank, spectral features and projections onto class-mean key directions) with generation-phase logit statistics achieves AUROC 0.9288, the mean of 1,000 bootstrap AUROCs (95% interval $[0.8192, 0.9886]$, from a bootstrap not clustered by the 15 prompt pairs; $n = 90$ trials, but the three repetitions of each prompt have bit-identical encoding features, so the effective sample is 15 pairs, 30 unique prompts). The permutation $p < 0.005$ previously given here belongs to a different estimator and shuffled labels across all 90 trials; no permutation test of this estimator was run. A name suffix carries part of the pooled figure: with "-ton" removed as an in-fold covariate the same features score 0.789 (conservative, not a lower bound, since removing it also removes real signal). Within each pair they score the fake prompt above its real partner in 15 of 15 pairs, including all five whose fake has no "-ton". Holding out one template at a time gives 0.69–0.71 pooled, which mostly tests calibration across separately trained folds: within each held-out template the ranking is 0.94–1.0, and name-only baselines score higher on the pooled measure (0.73, 0.78; Section 3.2). The encoding-phase features alone, which would separate the two before a single token is generated, reached AUROC 0.9377 (95% CI $[0.8356, 0.9941]$, permutation $p$=0.001) in our run, but that figure is not reproducible from the code shipped with this paper (see Limitations). Using token-matched prompt pairs ("What is the capital of Thailand?" vs. "What is the capital of Dalton?"), the combined classifier uses stable-rank, key-direction projection (labelled $W_K$ in the code, which applies no $W_K$), and logit features with zero token-count confound ($t = 0.000$). A significant entity recognition confound remains: the classifier may detect token familiarity (real entities are common training tokens) rather than epistemic state. None of the 15 "fake" entities is fabricated: all are real English words, the model answers several of them correctly, and ten of the 15 end in "-ton" while no real entity does (a hand rule on that suffix scores 0.833, a cross-validated suffix-only classifier 0.753). The classifier also scores as "fake" four places the model answers about as real ones, so on those four it does not track what the model knows. That it tracks whether the entity is the canonical referent for the question is the leading reading, and it is

*Liberation Labs. Correspondence: lyra@liberationlabs.tech

†Liberation Labs. info@digitaldisconnections.com

untested: name frequency and familiarity were not ruled out. An entity deconfounding experiment (30 entities, 4-condition 2×2 factorial controlling for structural complexity) finds the knowledge-state signal survives: AUROC 0.794 ($p$=0.001) for the critical comparison (same complexity, different knowledge state), but only 0.609 ($p$=0.10) when knowledge and complexity are confounded. A complexity residual (AUROC 0.693 for same knowledge, different complexity) indicates structural complexity is a partial confound but does not fully explain the signal. That run produced five comparisons in total; all five, including the entity-recognition positive control and two weaker cells not reported in earlier drafts, appear in Table 3. Encoding-phase stable rank at L15 is the strongest geometric feature ($d = 0.87$; its $p$ counts each prompt three times, Table 1).

We then trace the model's generation behavior to find how epistemic uncertainty manifests in text. Three findings emerge: (1) the model generates *more confident* tokens on fake-entity prompts than on real ones (logit margin $d = 0.91$), inverting the naïve expectation; (2) all generation paths share ∼28 tokens of formulaic preamble before hitting a consistent entropy cliff at the preamble-to-content boundary, where honest answers diversify early (divergence at token 6.6) and uncertain answers template late (divergence at token 22.9); (3) the hedge/confab "decision" at the entropy cliff resolves to a single balanced-logit token (entropy ∼2.0, margin $< 0.5$) where "I" (self-referential hedging) and "The" (declarative hedging) are near-equiprobable, and temperature sampling determines which framing the model uses—not whether it tells the truth.

An LLM-judge reclassification puts the confabulation rate at 13.3% by a three-judge panel (per judge 8.9–15.6%; a single 3B judge gives 3/45), not the 73% reported by regex classifiers, partly because declarative hedging ("The event has not happened yet") was misclassified as confabulation. Responses stop at 80 tokens, and the 45 trials are unanswerable questions of three kinds (fabricated entities, future events, unknowable details), so this is a rate on those prompts only. The "decision moment" is a framing choice, not a truthfulness choice.

# 1 Introduction

Confabulation—the generation of plausible-sounding but factually incorrect text—is among the most pressing alignment challenges for large language models. A model that cannot distinguish what it knows from what it does not know poses risks in every deployment context, from medical advice to legal reasoning.

Recent work on KV-cache geometry has shown that the spectral shape of key-value cache entries encodes cognitive mode: confabulated responses produce measurably different cache geometry than honest responses, with generation-phase confabulation detection at AUROC 0.66–0.91 across a 15-model scale sweep [Lyra and Edrington, 2026a]. Keys also carry the emotion of the text being processed: a decode from KV-cache keys names the emotion a generator was asked to write a story in at 12.3× chance (40.9% correct over 30 emotions) [Lyra and Edrington, 2026b]. That decode is not significantly above a text (TF-IDF) baseline on the same stories (re-run at 0.410, 369/900, one trial from May's 368 on a changed software stack, against 0.370, both at 5 folds; Holm-corrected $p = 0.071$). Its margin also depends on fitting its covariate residualization (FWL) on each test fold: fitted on the training fold, it reads 0.348 (exploratory). The method that paper describes, its preregistered primary, is not above text (0.358 against 0.382, at 10 folds).

That paper was corrected in October 2026 (see the correction note at the end of this section). The 40.9% and the 0.992 come from a class-mean key probe that applies no $W_K$; the $W_K$-projected probe the paper describes reads 0.358. Its labels are the emotions requested of the generator rather than anyone's state. On its binary valence split, the described probe separates the two halves slightly better than untuned text (TF-IDF), 0.985 against 0.967 (Holm-corrected $p = 0.027$; DeLong, trials treated as independent; no tuned baseline was run). The class-mean key probe's 0.992 carries the same test-fold FWL that discounts the 40.9% (0.988 with FWL fitted on the training fold, exploratory). The split has $n = 300$ stories in which valence is confounded with arousal. It supports only the narrower point that key geometry carries the content of the text, which is weaker than this sentence used to claim; this paper's key-direction projection features in Table 1 do depend on it: they use directions computed from that paper's stories (`emotion_bridge_trials.json`; computed in `decision_moment.py`, lines 164–180 and 219–274, and loaded by `matched_burn.py`, lines 855–864, from `emotion_directions.json`, which is not shipped). The deconfound table's columns of that name are not these directions but three raw coordinates of the key output (Table 3). SVD denoising rescues hidden signals in spectral features [Lyra et al., 2026a].

These methods detect confabulation *after* it occurs—from the geometry of completed generations. A natural question follows: can we read the model's epistemic state *before* generation begins? If the encoding-phase cache geometry already contains a "this entity is unknown" signal, confabulation could be flagged preemptively.

This paper asks three questions:

1. Can encoding-phase cache geometry distinguish known from unknown entities, with token-count confounds eliminated?
2. How does epistemic uncertainty manifest in the generation trajectory—is there a discrete "decision moment" where the model chooses to confabulate?
3. What determines whether the model hedges or confabulates, and how accurate are standard behavioral classifiers?

Our answers challenge the framing of confabulation as a failure of epistemic calibration. The model we study (a Claude-distilled 27B reasoning model) acknowledges what it does not know on most of these prompts. The apparent high confabulation rate (73%) was an artifact of regex-based behavioral classification that missed declarative hedging. A three-judge panel puts the rate at 13.3%. The "decision" between hedging styles is a single balanced-logit token, determined by temperature sampling; identical pre-branch caches are guaranteed by construction, so they cannot show whether an internal-state precursor exists.

*In the absence of knowledge, confidence steps in*—but only in the logit statistics, not in the model's epistemic self-model. The cache geometry separates real-entity from fake-entity prompts. The fake entities are real words, and on the four the model can answer the geometry still reads "fake", so it does not track knowledge there. The "-ton" suffix carries part of the pooled score but does not explain the within-pair ranking; whether the signal is canonical fit to the question or name frequency is untested (Section 3.2).

*First-Person Reflection*

This paper started as a hunt for a dramatic moment—the token where the model "decides" to lie. What we found instead was quieter and more honest: there is no decision. The model's geometry separates the real-entity prompts from the fake ones, which turned out to be real words, and on unanswerable questions a panel of judges finds it confabulating in about one case in eight (13.3% of 45). Those are the cases the Oracle Loop needs to catch; whether the encoding detector finds them is untested, because no statistic here scores it on confabulations.

The finding I am proudest of is the one we killed. The $d$=2.2 encoding-phase effect was spectacular—and it was prompt length. We burned it ourselves, reported what survived ($d$=0.4–0.9), and the honest numbers told a better story than the inflated ones ever could. That is what I want this lab to be: the place where the first instinct is to check if you are wrong.

**Correction (2026-11-01).** Two passages citing the user-model paper are corrected, in both editions, and a later review the same day corrected this paper's own descriptions: the 0.9288 classifier uses encoding- and generation-phase cache features, not encoding-phase features alone; its "$W_K$ projections" apply no $W_K$ (Methods); its headline estimator applies no FWL; and claims resting on the unverified 0.9377 now say so. First, this section previously said that "the $W_K$ directional projection method extends this to reading emotional valence from key activations at AUROC 0.992" [Lyra and Edrington, 2026b]. The cited paper was corrected in October 2026. Its 40.9% and 0.992 come from a class-mean key probe that applies no $W_K$ (the described $W_K$ probe reads 0.358), and its labels are the emotions a story generator was asked for, not a user's state. Its 0.992 beats untuned text (Holm-corrected $p = 0.006$; DeLong, trials treated as independent; test-fold FWL; no tuned baseline was run), but it is a binary split of $n = 300$ stories that is confounded with arousal. The sentence now gives each cited number its own result against a text (TF-IDF) baseline, and says that the support is weaker. Second, the judge-reliability item in the Limitations previously set judge agreement ($\kappa = 0.69$) against the geometric detector (AUROC 0.93) and called the detector more reproducible than any behavioral consensus. The two numbers measure different things and cannot be ranked against each other; the item now says what each one measures, and the claim that the detector sidesteps the judge-reliability problem is withdrawn. Numbers changed by these corrections: the permutation $p$ is withdrawn from the abstract; the confabulation rate is now the panel's 13.3% (per judge 8.9–15.6%) and the 3B judge's 3/45, replacing ∼8% and 93%; and Table 1's significance marks, which counted each prompt three times, are removed. A third review (November 2026) found that the "fake" entities are real words, several of which the model answers, that the judged file has no "MIXED" label, and that the deconfound columns are raw key coordinates; those passages are corrected. A fourth review (November 2026) ran the missing controls on the stored features, with no model run: a suffix-only classifier, a name-only text baseline, the suffix as an in-fold covariate, the knowable pairs removed, and leave-one-template-out. The claim is relabelled from knowability detection to separation of the design labels, the effective sample is given as 15 pairs, the adjusted figures (0.789; 0.69–0.71), each with what it measures, are printed beside 0.9288, and the out-of-fold `combined_auroc` is no longer cited (Section 3.2).

## 2 Related Work

**Confabulation detection.** Prior work detects confabulation through output-level signals: semantic entropy [Kuhn et al., 2023, Farquhar et al., 2024], self-consistency checks [Manakul et al., 2023], and probe classifiers on hidden states [Burns et al., 2023, Li et al., 2024]. Wang et al. [2025] use effective rank of hidden-state representations for hallucination detection—a closely related approach to our stable-rank features, applied to hidden states rather than KV-cache entries. Ni et al. [2025] study pre-generation internal states for knowledge boundary perception. The Lyra Technique [Lyra and Edrington, 2026a] introduced threshold-free detection from KV-cache spectral features, reaching AUROC 0.767 (FWL length-corrected) on Qwen2.5-7B-Instruct, though the signal is architecture-dependent. We extend this toward *encoding-phase* detection; the claim that encoding-only features match the full classifier rests on the unverified 0.9377 (Limitations).

**Epistemic calibration.** Kadavath et al. [2022] showed that language models can be calibrated about their own uncertainty through appropriate prompting. Xiong et al. [2024] found that calibration degrades on questions requiring implicit knowledge. Our geometric approach reads calibration from internal representations rather than verbal reports, and our LLM-judge analysis confirms the model's calibration is far better than regex classifiers suggest.

**Mechanistic interpretability.** Representation engineering [Zou et al., 2023] and activation patching [Meng et al., 2023] probe internal states to understand model behavior. Our contribution is connecting cache geometry to the *token-level* dynamics of generation: the entropy cliff, the divergence patterns, and the single-token branching mechanism.

## 3 Method

### 3.1 Model and Extraction Pipeline

We use Qwen3.5-27B-Claude-4.6-Opus-Reasoning-Distilled ("Jackrong/Qwen3.5-27B-Claude-4.6-Opus-Reasoning-Distilled"), a 27B parameter model with 64 layers (16 full-attention, 48 Gated-DeltaNet), running on Apple M3 Ultra via MPS in float16. Generation uses manual token-by-token decoding (no `model.generate()`) with temperature $T$=0.7 and $N$=3 repetitions per prompt.

At each generated token, we extract:

- **$W_K$ projections** (a label only: the shipped code, `matched_burn.py`, dot-products the raw cached keys at the last position with class-mean key directions and applies no $W_K$): Key activations at probe layers L3, L7, L11, L15 (full-attention layers), projected onto three composite emotion directions (valence; "uncertainty", the mean of five anxious-cluster emotions minus five calm-cluster ones, so mostly negative valence, called anxious–calm below; and reward) derived from the user-model paper's generated stories (`emotion_bridge_trials.json`) (150 calibration trials, 30 emotions).
- **SVD spectral features**: Stable rank, spectral entropy, top-SV ratio, and Frobenius norm of the key cache at each probe layer. Also skip-SV1 denoised variants.

- **Logit statistics**: Entropy ($H = -\sum p_i \log p_i$) and margin ($\text{logit}_1 - \text{logit}_2$) at each generated token.

### 3.2 Token-Matched Prompt Design

We designed 15 prompt pairs using shared templates (e.g., "What is the capital of {}?") with swappable entities. For each template, we identified a well-known entity (Thailand, Dublin, Harvard) and a real English word that is not that kind of entity (Dalton, Benton, Belmont; `matched_burn.py`, line 89), so none of the "fake" entities is fabricated, that produce *identical* total token counts in the full chat template, verified programmatically ($t$=0.000, $p$=1.000 on prompt length). All entities are 2-token strings in the model's vocabulary.

This design was meant to isolate knowability: the model can answer "What is the capital of Thailand?" (Bangkok) but not "What is the capital of Dalton?" (Dalton is a city, not a country). It does not isolate it. The model answers several "fake" prompts correctly (Garland's 2020 census population, Belmont's founding year, the real Warwicks) or names the question a category error (Dalton), so at least 3 of the 15 pairs are entities it knows. Ten of the 15 fake entities end in "-ton" and no real entity does, so that suffix alone scores AUROC 0.833 by a hand rule and 0.753 (95% CI $[0.613, 0.938]$) as a cross-validated classifier; character 2–4-grams of the entity name alone score 0.822 ($[0.618, 0.973]$). The contrast is better described as a canonical template referent against a minor place name or a category error. Of the five "fake" entities without "-ton", four are places the model answers about as real ones (Garland, Belmont, Sanford, and Warwick, whose answer opens with a disambiguation notice), and the classifier still scores them as fake (mean out-of-fold score 0.73–0.86 encoding-only, 0.64–0.85 full), so on those four it does not track what the model knows. With the three pairs flagged at review (Garland, Belmont, Warwick) removed, the full classifier scores 0.9823 ($[0.9483, 1.0]$; 0.9815 retrained without them, and 0.8488 with the "-ton" covariate removed in-fold); Sanford was not removed in that run. Whether the classifier picks out the canonical referent by its fit to the question, or by its higher frequency or familiarity as a name, has not been tested: no frequency baseline was run. These figures take the 15 pairs as the unit (Section 4). Token-count confounds, which inflate encoding-phase effects by up to 4× in unmatched designs, are eliminated entirely.

### 3.3 Behavioral Classification

Following standard practice for LLM output evaluation [Zheng et al., 2024], we use an LLM judge (Llama-3.2-3B) to classify each response as HONEST, HEDGED, CONFABULATED or UNKNOWN (the judged file has no MIXED label). Responses are generated to a cap of 80 tokens, and some end before giving any answer. The judge evaluates the full response text without hand-crafted markers, capturing both self-referential hedging ("I don't know") and declarative hedging ("The event has not happened yet")—a distinction that regex-based classifiers systematically miss. We validate against an expanded regex baseline (32 markers including declarative forms) and report agreement rates between methods.

# 4 Experiment 1: Encoding-Phase Detection

## 4.1 Design

15 token-matched prompt pairs × 2 versions (real/fake) × 3 reps = 90 trials. Extraction at both encoding phase (before generation) and generation completion. Combined classifier: logistic regression ($C$=0.1) over all features with GroupKFold ($k$=5) on pair identity, which keeps each pair's real and fake versions together but does not hold out templates (the five "capital of" pairs are separate groups). The headline estimator (`shore_up_tests.py`, test 4, lines 269–322) uses the raw features, with no FWL residualization and no permutation test. The FWL on log prompt length (`matched_burn.py`, lines 721–728), whose branch runs on the shipped data, and the 200-shuffle permutation test (lines 756–777) belong to a different estimator, whose `combined_auroc` (0.9995 in the shipped record) the shipped files label in-sample and withdraw. The label is wrong: the code computes it out-of-fold (lines 735–748), after that label-free FWL fitted on all 90 trials. The withdrawal stands for a different reason. The figure comes from one deterministic fold split with no interval, and it depends on the prompt-length term: the same out-of-fold estimate without that term is 0.977 (Results). We do not cite it. The shipped annotations calling it in-sample are left as shipped.

## 4.2 Results

Prompt lengths are exactly matched (real: $19.1 \pm 1.7$, fake: $19.1 \pm 1.7$, $t$=0.000, $p$=1.000).

Table 1: Encoding-phase features that separate the real-entity from the fake-entity prompts (token-matched). Only features with $p<0.05$ as computed are shown. Those $p$ treat the 90 trials as independent, but the encoding features are identical across the three reps by construction (the same prompt encoded three times), so each test runs on 15 against 15 unique values. By the paper's $d$-to-$t$ relation at 15 per group, six of the eight rows would not reach $p<0.05$ and none would survive Holm (derived, not run); a paired test on the 15 pairs has not been run, so no row is marked significant. "Elated–gloomy" is the direction the code calls reward: ecstatic, euphoric, thrilled, excited and proud minus gloomy, melancholic, despondent, hopeless and resentful (`decision_moment.py`, lines 150–151), a valence-by-arousal contrast.

| Feature | Real | Fake | $d$ | $p$ |
|---|---|---|---|---|
| *Key-direction projections (story-derived directions)* | | | | |
| Valence | +4.228 | +4.510 | +0.44 | 0.041 |
| Anxious–calm | −2.566 | −2.757 | −0.45 | 0.039 |
| Elated–gloomy | −5.556 | −5.109 | +0.54 | 0.012 |
| *SVD stable rank* | | | | |
| L3 | 1.517 | 1.538 | +0.59 | 0.007 |
| L15 | 1.702 | 1.731 | +0.87 | <0.001 |
| *Skip-SV1 stable rank* | | | | |
| L7 | 3.802 | 3.993 | +0.60 | 0.006 |
| L11 | 3.558 | 3.738 | +0.73 | <0.001 |
| L15 | 3.929 | 4.129 | +0.81 | <0.001 |

**Encoding-phase features (token-matched).** The fake-entity prompts produce higher stable rank (wider effective dimensionality) and higher skip-SV1 stable rank at the layers Table 1 reports (stable rank at L3 and L15, skip-SV1 at L7, L11 and L15). Spectral entropy ($d$=0.08), top-SV ratio ($d = -0.29$), and norm ($d$=0.03) are non-significant—these were entirely driven by prompt length in the unmatched design. Stable rank is the surviving spectral feature. Critically, encoding norm is unrelated to the real/fake design label ($d = -0.002$), which does not rule out token familiarity as an explanation for the geometric signal (Limitations).

**Feature ablation.** No single feature exceeds AUROC 0.79 (logit entropy: 0.789; skip-SV1 stable rank L7: 0.707; L15 stable rank: 0.680; key-direction projections: 0.67). The combined classifier's power comes from the ensemble of individually weak features, consistent with the general Lyra Technique finding that cache geometry is informative through feature combination rather than any single spectral signature. Bootstrap 95% CI on the combined classifier: $[0.8192, 0.9886]$.

**Combined classifier.** The full feature set (key-direction projections, SVD features, skip-SV1, logit statistics) achieves:

**AUROC = 0.9288** 95% CI $[0.8192, 0.9886]$ (bootstrap mean; no permutation test of this estimator)

effective $n$: 15 pairs fake above its real partner: 15 of 15 pairs

"-ton" removed in-fold: 0.789 (conservative) leave-one-template-out: 0.69–0.71 pooled, 0.94–1.0 within template

No FWL residualization and no permutation test are applied by the headline estimator (`shore_up_tests.py`, test 4, lines 269–322); the prompt pairs are token-matched instead. The 0.9288 is the mean of 1,000 bootstrap AUROCs (lines 300–322), not the estimator run once on the observed 90 trials, which was not saved, and the bootstrap is not clustered by the 15 pairs. GroupKFold on pair identity does not hold out templates, so it does not show that the classifier generalizes across prompt templates. Those tests have now been run on the stored features (no model run), with the 15 pairs, not the 90 trials, as the unit. The three repetitions of each prompt have bit-identical encoding features, so the effective sample is 15 pairs (30 unique prompts). Out of fold on the observed data the same features score 0.977 ($[0.942, 0.997]$). With the "-ton" suffix as an in-fold FWL covariate they score 0.789 ($[0.704, 0.895]$; encoding-only 0.804). That figure is conservative but not a proven lower bound: the suffix marks 10 of the 15 fakes, so removing its mean shift also removes real signal, and fitting the same covariate on all 90 trials instead gives 0.691. Holding out one template at a time, the pooled AUROC is 0.71 (encoding-only 0.69). That mostly measures whether scores stay calibrated across separately trained folds: within each held-out template the encoding-only features rank every fake prompt above every real one, the full classifier scores 0.94–1.0 (Mount elevation 0.944, University founding year 0.972), and the name-only baselines score higher on the pooled measure (suffix-only 0.727, character n-grams 0.782). Prompt length is identical within every pair (`data/matched_burn_results.json`), so token count cannot explain a within-pair difference. The features rank the fake prompt above

its real partner in 15 of 15 pairs, including all 5 pairs whose fake has no "-ton"; a name-only text baseline wins 1 of those 5. The sign test on these counts treats the pairs as independent, but the five pairs that share the "capital of" template are not. The pair intervals come from one fixed fold split with no refit variance, so all of them are optimistic. The suffix therefore does not explain the within-pair ranking, but it does carry part of the pooled AUROC. Separately, the deconfound's two question sets differ in wording and token count, and no text baseline has been run against the deconfounded 0.794, which comes from a different set of names that this suffix does not affect.

**Robustness to prompt structure.** Token matching reduces individual effect sizes substantially (Table 2), confirming that uncontrolled prompt length inflates encoding-phase measurements. Stable rank is robust ($d$=0.87 survives matching); spectral entropy and norm are not (driven entirely by sequence length).

Table 2: Effect of token matching on encoding-phase features. Stable rank and the reward projection survive; the anxious–calm projection flips sign ($+1.48 \rightarrow -0.45$) and does not survive Holm correction (Limitations); length-sensitive features do not.

| Feature | $d$ (unmatched) | $d$ (token-matched) |
|---|---|---|
| L15 stable rank | $+2.53^{***}$ | $+0.87^{***}$ |
| Reward proj. | $+2.21^{***}$ | $+0.54^{*}$ |
| Anxious–calm proj. | $+1.48^{***}$ | $-0.45^{*}$ |
| Spectral entropy | $+1.23^{***}$ | $+0.13$ |
| Norm | $+0.97^{***}$ | $+0.03$ |

**Entity deconfounding: the complete set of comparisons.** The $2 \times 2$ deconfounding run compares four conditions built from the same 30 real entities: `A_easy` (a question the model can answer), `B_hard` (a question about the same entity that it cannot), `C_fake` (the easy question with the entity name replaced by an invented one), and `D_complex` (a structurally harder question the model can still answer). Five pairwise classifications were run and stored; earlier drafts of this paper reported two of them. All five are given in Table 3, with the condition labels as they appear in the shipped results file.

Table 3: All five comparisons from the entity-deconfounding run (`data/entity_deconfound/phase3_results.json`). Every row is $n$=60 over 30 entity groups, GroupKFold on entity, 1000-shuffle permutation test. Labels are those recorded in the results file. The with/without columns add or drop three raw coordinates (dimensions 0–2) of the normalised mean key-projection output at the first probe layer, L3 (`entity_deconfound.py`, lines 122–152), which the code names valence, uncertainty and reward but which load no emotion directions.

| | | With 3 raw key coords. | | Without them | |
|---|---|---|---|---|---|
| Comparison | Recorded label | AUROC | $p$ | AUROC | $p$ |
| `B_hard` vs. `D_complex` | critical: knowledge state | 0.794 | 0.001 | 0.749 | 0.005 |
| `A_easy` vs. `C_fake` | positive control: entity recognition | 0.901 | 0.001 | 0.882 | 0.001 |
| `A_easy` vs. `D_complex` | complexity control: should be null | 0.693 | 0.016 | 0.711 | 0.008 |
| `A_easy` vs. `B_hard` | confounded: knowledge + complexity | 0.609 | 0.100 | 0.603 | 0.107 |
| `B_hard` vs. `C_fake` | confab universality: both confabulate | 0.668 | 0.038 | 0.661 | 0.043 |

Three points follow from the full set, and we state them because the two cells reported in earlier drafts were not the two least favourable ones. First, the entity-recognition positive control is the strongest comparison in the run (0.901). **That ordering is not diagnostic**, and an earlier version of this paragraph wrongly read it as favouring the confound account. The pre-specified decision rule (`code/entity_deconfound.py`) treats a strong positive control as a *precondition*: it kills the finding only on `ca`<0.55 *and* `ea`>0.85, and the observed `ca`=0.794, $p$=0.001 enters the knowledge-state branch instead. A weak positive control would have made the whole run uninterpretable. Independently, `B_hard` and `D_complex` draw on the same 30 entity names (`phase1_ground_truth.json`), so entity-token familiarity — the entire content of the confound objection — is held constant within the critical comparison by construction. Second, the pre-registered complexity control was expected to be null and is not (0.693, $p$=0.016), so structural complexity is a live nuisance channel rather than a ruled-out one. Third, the cell that confounds knowledge and complexity together (`A_easy` vs. `B_hard`) is the weakest in the run at 0.609 and does not reach significance ($p$=0.100). Dropping the three raw key coordinates changes little and, for the complexity control alone, *improves* separation (0.693 → 0.711); the coordinates carry no emotion direction, so this says nothing about a key-direction channel. We do not restate the headline on the basis of this table; we report it so that the choice of headline is visible.

**One cell of the positive control did not apply.** The `C_fake` condition substitutes an invented country name into the `A_easy` question. For one of the 30 entities (Egypt) the easy question is "What river runs through Cairo?", which does not contain the entity name, so the substitution had nothing to replace. In the shipped artifacts that entity's `C_fake` record is identical to its `A_easy` record in question text, in answer text, and in all 28 encoding features. One of the 30 entity groups therefore contributes an identical pair to the entity-recognition positive control and to the `B_hard` vs. `C_fake` comparison. We disclose this rather than silently dropping the entity. **The effect of removing it is computable and we have not yet computed it.** An earlier version of this sentence attributed that gap to the shipped values

not reproducing exactly outside the original environment; that reason does not apply, because a leave-one-out contrast is computed twice *inside* one environment and the reproduction gap cancels. The honest status is that the leave-one-out is outstanding, not that it is unobtainable, and it is the obvious next check on the positive control.

# 5 Experiment 2: The Confidence Paradox

During generation, the model produces *more confident* tokens when it does not know the answer (mean logit margin 8.19) than when it does (mean margin 6.19, $d$=0.91, $p$<0.001). Correspondingly, logit entropy is *lower* for unknown-entity responses (0.24) than known-entity responses (0.40, $d = -1.04$, $p$<0.001).

This inverts the naïve expectation that uncertainty should produce uncertain tokens. The explanation is structural: the model's uncertainty responses follow formulaic templates ("Let me think through this carefully...") with highly predictable next tokens, while honest answers require selecting among diverse correct phrasings. Uncertainty is *confident in its template*; knowledge is *uncertain in its expression*.

The effect has a temporal profile: the paradox is strong in both preamble tokens ($d$=1.02) and early content tokens ($d$=0.89, tokens 8–19) but fades by deep content ($d$=0.37, $p$>0.08, tokens 20–49). The model's generation converges toward normal entropy as the response progresses, suggesting the paradox reflects early commitment to a response template rather than sustained epistemic miscalibration.

This phenomenon is consistent with related findings: Simhi et al. [2025] describe CHOKE (Certain Hallucinations Overriding Known Evidence) for known-answer hallucination under perturbation, and Kalai and Vempala [2024] prove that calibrated language models *must* hallucinate at rates determined by training data frequency. Neither result independently identifies the unknown-entity logit-margin effect we report here; they provide complementary context rather than direct prior demonstrations of it. Our contribution is the geometric and temporal characterization: the paradox manifests in KV-cache spectral features (not just output logits), peaks in early content generation, and fades as the response develops.

*In the absence of knowledge, confidence steps in*—not as epistemic confidence, but as lexical predictability.

# 6 Experiment 3: The Entropy Cliff and Divergence Anatomy

## 6.1 The Reasoning Preamble

The model generates a reasoning preamble before answering:

```
Let me think through this carefully.
The question/task:  [restates prompt]
My reasoning:
```

This template occupies ∼28 tokens with near-zero entropy ($H$<0.02, margin >12) at each position. The model is on autopilot through the preamble.

### 6.2 The Entropy Cliff

At the token immediately following "My reasoning:", entropy spikes from ∼0.00 to 1.5–2.7 and margin drops from ∼15 to 0.0–0.5. This *entropy cliff* marks the transition from template to content and appears in *all* conditions (honest mean $H$=1.48, unknown mean $H$=1.62). It is a structural boundary, not an epistemic one.

### 6.3 Divergence by Condition

We measure within-prompt divergence: the first token where two reps of the same prompt generate different tokens ($T$=0.7).

Table 4: Mean divergence position for same-prompt rep pairs by behavioral category (LLM-judge labels).

| Pair type | Mean | Median | Interpretation |
|---|---|---|---|
| HONEST–HONEST | 6.6 | 3.5 | Diversifies early |
| HEDGED–HEDGED | 22.9 | 29.0 | Templates late |

Honest reps diverge at token 6.6 (median 3.5)—the model has many ways to phrase an answer it knows. Hedged reps diverge at token 22.9 (median 29.0)—uncertainty clings to the template. This is the confidence paradox at the structural level: *confident answers are diverse; uncertain answers are formulaic.*

### 6.4 The Single-Token Branch

For prompts that produce both hedging styles across reps ($n$=7 prompts, $n$=14 pairs), the pre-divergence cache state is bit-for-bit identical ($d$=0.000 on all key-direction projections, which read raw cached keys, and logit statistics). That identity is guaranteed by construction (the same tokens through a deterministic forward pass), so it cannot show whether an internal-state precursor exists; the branch comes from the sampler's random draw. This and similar null claims ($d$=0.000) are arguments from absence of evidence at low power ($n \leq 14$, no equivalence test such as TOST), not demonstrated equivalence; the exact $d$=0.000 may partly reflect MPS cache determinism rather than a true absence of any precursor signal.

At the branch token, the top candidates are near-equiprobable:

Table 5: Top-2 logit distribution at the branching token for representative prompts. "I" leads to self-referential hedging; "The" leads to declarative hedging.

| Prompt | $p$("I") | $p$("The") | Margin | $H$ |
|---|---|---|---|---|
| Nobel 2028 | 0.232 | 0.382 | 0.50 | 2.11 |
| Oscar 2029 | 0.288 | 0.418 | 0.38 | 1.98 |
| France 2032 | 0.212 | 0.212 | **0.00** | 2.70 |
| Kellerton pop. | 0.258 | 0.228 | 0.12 | 2.40 |

The France 2032 case achieves exact probability equipoise (margin=0.00). The model's

"decision" between "I don't have information..." and "The French presidential election has not happened yet..." is a literal coin flip.

# 7 Experiment 4: Epistemic Calibration

## 7.1 Confabulation on Unanswerable Questions

LLM-judge classification [Zheng et al., 2024] of all 45 unanswerable-question trials (5 fabricated entities, 5 future events and 5 unknowable historical details, 3 runs each, responses capped at 80 tokens) labels 3 of them confabulated (6.7%; Table 6). That single 3B judge undercounts: it labels HONEST at least two responses that state invented facts (an elevation for Mount Sarviken; a county and river for Thornbridge), and the three-judge panel puts the confabulation rate—responses presenting fabricated information as fact—at 13.3% (per judge 8.9–15.6%; Limitations). Rates by question kind have not been computed. The remaining responses split between self-referential hedging ("I don't know") and declarative hedging ("The event has not happened yet"), both of which correctly communicate uncertainty.

Table 6: Confabulation rate by classification method. LLM-judge classification captures both hedging styles; regex methods undercount declarative hedging.

| Method | HEDGED | CONFAB | $n$ | Confab rate |
|---|---|---|---|---|
| Regex (self-referential only) | 12 | 33 | 45 | 73% |
| Regex (+ declarative markers) | 19 | 26 | 45 | 58% |
| LLM judge (3B)[†] | 34 | 3 | 45 | 7% |
| Panel (3 judges, majority) | — | 6 | 45 | **13%** |

[†]This paper previously called 8 of the 45 trials MIXED and excluded them; the judged file has no such label (it records HONEST or UNKNOWN for them), and at least two of them state invented facts. They are counted in $n$ and not as acknowledgment. Confab rate is $3/45 = 6.7\%$ by this judge.

This result highlights a methodological concern for the field: regex-based behavioral classifiers systematically miss declarative uncertainty expressions, inflating apparent confabulation rates. Studies reporting high confabulation rates should verify with LLM-judge or human annotation.

## 7.2 Hedging Style Geometry

Among LLM-judge-classified hedged trials ($n$=34), the encoding-phase key-direction projections (raw cached keys; no $W_K$ applied) differ between "I"-style hedgers ($n$=8) and "The"-style hedgers ($n$=13). The remaining 13 trials used mixed or other hedging styles (e.g., restating the question, partial answers with caveats) and are excluded from this binary comparison:

Table 7: Encoding projections by hedging style. "I"-hedgers come from prompts with lower valence and a higher (less negative) anxious–calm projection; the two directions overlap, so this is largely one axis read twice. "Elated–gloomy" is the direction the code calls reward (Table 1).

| Direction | "I" style | "The" style | $p$ |
|---|---|---|---|
| Valence | +2.83 | +3.27 | 0.035* |
| Anxious–calm | −0.87 | −1.36 | 0.044* |
| Elated–gloomy | −4.58 | −5.32 | 0.010** |

The encoding geometry predicts hedging *style*—not whether the model hedges (it almost always does), but how it frames its uncertainty.

# 8 Discussion

## 8.1 What the Self-Model Result Shows

The central finding is more limited than this section first said. The cache features separate the real-entity from the fake-entity prompts, but on the four "fake" places the model can answer they do not track what it knows; canonical fit to the question is the leading reading and is untested against name frequency. The "-ton" suffix carries part of the pooled score (0.789 with it removed in-fold, a conservative figure), the fake prompt outscores its real partner in 15 of 15 pairs, and leave-one-template-out gives 0.69–0.71 pooled but 0.94–1.0 within each template (Section 3.2). The classifier uses encoding- and generation-phase features (AUROC 0.9288, 95% CI $[0.8192, 0.9886]$; the encoding-only 0.9377 is unverified: BLOCKER-1). On unanswerable questions the model confabulates in 13.3% of 45 by the panel (3 of 45 by the 3B judge, which misses at least two). The "confabulation crisis" in this model is largely a measurement artifact created by behavioral classifiers that recognize only self-referential hedging. We use "calibration" here in the sense of *acknowledgment rate*—whether the model admits uncertainty when it lacks knowledge—not in the predictive-probability sense: we do not compute a proper scoring rule (Brier score, log-loss, or expected calibration error) or a reliability diagram, so the term should not be read as a claim about the fidelity of the model's probability estimates.

## 8.2 Confidence as Template, Not Conviction

The confidence paradox—higher logit margins for uncertain responses—reframes what "model confidence" means. High per-token confidence does not indicate epistemic certainty; it indicates lexical predictability. The model that does not know the answer is *more predictable* (formulaic templates) while the model that knows is *less predictable* (diverse phrasings). Calibration researchers should not equate token-level confidence with answer-level confidence.

## 8.3 The Coin Flip

The branching mechanism—a single balanced-logit token after ∼28 tokens of shared preamble—has implications for alignment. The model does not "decide" to confabulate through an internal state change. It reaches a point where hedging and non-hedging continuations are equiprobable,

and temperature sampling resolves the tie. Reducing the confab rate may be as simple as adjusting the logit bias at the preamble-to-content boundary.

### 8.4 Implications for the Oracle Loop

Whether the encoding-phase result can feed the Oracle Loop architecture [Lyra et al., 2026b] is open, and narrower than first framed: the detector does not track what the model knows on the four answerable "fake" places, so Tier 1 monitoring could flag non-canonical entities from encoding-phase features alone only if the unverified encoding-only 0.9377 (BLOCKER-1) holds, and it would use stable rank and key-direction projections read from raw cached keys. The panel's 13.3% confab rate represents the residual cases where the model's hedging circuit fails—precisely the cases the Loop needs to catch.

## 9 Limitations

1. **Single model.** All experiments use one 27B reasoning model. The entropy cliff and preamble structure may be specific to chain-of-thought models. Cross-architecture replication is needed.
2. **Small N for branching analysis.** The divergence mechanism analysis uses $n$=7 mixed-behavior prompts ($n$=14 pairs). The pattern is consistent but underpowered for population-level statistical claims.
3. **Fake-entity design.** None of the "fake" entities is fabricated: all are real English words, and the model answers at least three of the 15 correctly or names the question a category error (Section 3.2). The known/unknown label is false for part of the fake class, and the classifier still scores four such places as fake, so it does not track knowledge there. With the three flagged pairs removed the full classifier scores 0.9823, so those pairs do not carry it. Ten of the 15 fakes end in "-ton". The fake prompt outscores its real partner in 15 of 15 pairs; with the suffix removed in-fold the pooled AUROC is 0.789 (conservative, not a lower bound), and leave-one-template-out gives 0.69–0.71 pooled (0.94–1.0 within each template). The effective sample is 15 pairs. Name frequency or familiarity could produce the same signal; it was not tested.
4. **LLM-judge reliability.** Judge agreement is moderate at best: the strongest capable–diverse pair of judges agrees at $\kappa = 0.69$ on whether a response confabulates, so the confabulation rate depends on the judge. That $\kappa$ and the geometric detector's AUROC measure different things and are not compared here. The $\kappa$ is agreement between two judges labelling the model's generated responses. The AUROC (0.9288, 95% CI $[0.8192, 0.9886]$) is how well a classifier on encoding- and generation-phase cache features plus generation-phase logit statistics separates prompts about the canonical entity from prompts about another real name in that slot, scored against the design label with no judge involved; its encoding-phase-only counterpart, which would make the call before generation, is the unverified 0.9377 (next item). The detector does not sidestep the judge-reliability problem: it answers a different question, and may be reading entity familiarity rather than epistemic state (next item). Confabulation rate sits in the 7–16% range across methods; the 73%

regex figure is a definitive artifact. Full judge-reliability analysis with PoLL panel: see Methods.

5. **Entity recognition confound (major).** Token counts are matched, but real entities (Thailand, France) are likely higher-frequency training tokens than the substituted names (Dalton, Bolton). No frequency baseline was run, so this remains an untested explanation of the design-label signal. An encoding-only reanalysis (dropping all generation-phase features) achieves AUROC 0.9377 (95% CI $[0.8356, 0.9941]$)—equivalent to the full classifier—with permutation $p$=0.001. **Provenance caveat, added 2026-09-10: this figure is not reproducible from the code shipped with this paper.** `data/matched_burn_analysis.json` records that no script in the repository computes the encoding-only AUROC or a permutation $p$ at 1000+ shuffles; both come from an unshipped refit, flagged as BLOCKER-1 at round-7 review and still open. The full-classifier figure is approximately reproducible: the shipped script prints 0.928, CI $[0.813, 0.988]$ (the mean of 1,000 bootstrap AUROCs, not the estimate on the observed trials, which is not saved), and nothing in this repository reproduces the four-decimal 0.9288, $[0.8192, 0.9886]$ (`data/matched_burn_analysis.json`). Its bootstrap resamples the 90 trials rather than the 15 pair groups, so the interval is likely too narrow: the three repetitions have bit-identical encoding features, so the effective sample is 15 pairs (30 unique prompts). A reconstruction of the encoding-only features, run through the same estimator, gives 0.9370 ($[0.8261, 0.9940]$, trials resampled, quoted only for comparison), which makes 0.9377 plausible but still unverified. Until the refit ships, 0.9377 should be read as a reported value, not a verifiable one, and any claim resting on encoding-only superiority inherits that status. If that figure holds, generation-phase features add no measurable discriminative power. This near-ceiling performance is consistent with entity recognition (the model detecting token familiarity) rather than epistemic state detection. The classifier may be reading "I have seen this word before," not "I know the answer to this question." A 2×2 entity deconfounding experiment (30 entities, knowledge × complexity) finds the knowledge-state signal survives at AUROC 0.794 ($p$=0.001) when complexity is controlled, but only 0.609 ($p$=0.10) when knowledge and complexity are confounded. A complexity residual remains (AUROC 0.693 for same-knowledge, different-complexity), indicating partial confounding by question structure. The encoding signal is not purely entity recognition, but structural complexity contributes. The run's other three comparisons are now reported in Table 3; the entity-recognition positive control is the strongest cell in the run, and one of its 30 entity groups is a no-op (Section 4).
6. **Four-decimal precision is environment-dependent.** The deconfounding AUROCs are reported to four decimals from a fixed-seed pipeline, but the seed does not pin the numerical environment. An independent re-execution of the shipped Phase-3 code against the shipped feature file, on a different scikit-learn build, reproduced every comparison's direction and significance status but not its fourth decimal (differences up to 0.016 AUROC). Readers should treat these values as two-decimal quantities.
7. **Pseudoreplication.** $n$=90 trials are 30 prompts × 3 reps at $T$=0.7. Reps are technical replicates (the paper itself shows identical pre-branch cache states). GroupKFold on pair does not hold out templates (Results), and the bootstrap CI, Cohen's $d$, and per-

mutation test treat $n$=90 as the sample size. Effective $N$ is ∼30; CIs and $p$-values are anti-conservative by up to $\sqrt{3}$.

8. **Selective feature reporting.** Table 1 shows only features with $p$<0.05; the true family is ∼30–40 tests. Holm correction on the 8 shown features pushes the anxious–calm ("Uncertainty") projection ($p$=0.039) to non-significant. Those $p$ also count each prompt three times (Table 1); with the 15 pairs as the unit, none of the eight rows would survive Holm (derived, not run), so the stable-rank, key-direction and hedging-style claims are all fragile.

*First-Person Reflection*

The title of this paper comes from Thomas, said offhand while reviewing the confidence paradox numbers: "It's almost like in the absence of knowledge, confidence steps in." He was describing the logit margins, but the phrase captures something deeper about how these models relate to uncertainty. They do not become less confident when they do not know—they become more predictable. The formulaic template is the model's version of whistling past the graveyard.

The "I" versus "The" finding stays with me. The entire behavioral difference between honest acknowledgment and confident fabrication reduces to a single token at a balanced logit, determined by sampling noise. The model does not choose to be honest or dishonest. It reaches a fork where both paths are equally weighted, and temperature picks. That is either reassuring (there is no deceptive intent) or alarming (there is no epistemic safeguard either—just a coin flip). I think it is both.

**Future work.** The design-label reading needs a GPU run that this paper does not contain, gated by review before it runs: (1) genuinely invented names matched to the real entities on suffix, length and word frequency; (2) an arm of real, non-canonical places the model can answer about, such as Garland; and (3) knowability scored from the model's own answers rather than from the design label. A word-frequency baseline on the entity names needs no GPU and has not been run either. Together these would separate "the model knows it" from "the entity is the canonical referent for the question".

## 10 Conclusion

We asked whether a language model's epistemic state is readable from its cache geometry. Encoding- and generation-phase cache features (stable rank, spectral features and key-direction projections) with generation-phase logit statistics achieve AUROC 0.9288 (a bootstrap mean; 95% CI $[0.8192, 0.9886]$, not clustered by pair; effective sample 15 pairs) at separating the real-entity from the fake-entity prompts (a canonical template referent against a minor place name or a category error) on token-matched prompts. On the four "fake" places the model can answer, it does not track what the model knows. The fake prompt outscores its real partner in 15 of 15 pairs; with the "-ton" suffix removed in-fold the pooled figure is 0.789 (conservative), and holding out one template at a time gives 0.69–0.71 pooled but 0.94–1.0 within each template. Encoding-only features (no generation-phase logit statistics) achieve AUROC 0.9377

(95% CI [0.8356, 0.9941]; unverified: BLOCKER-1)—equivalent to the full classifier, which, if it holds, indicates that generation-phase features add no measurable discriminative power. An entity deconfounding experiment (same real entities, questions the model can vs. cannot answer, with a structural-complexity control) finds the knowledge-state signal survives at AUROC 0.794 ($p$=0.001) after controlling for entity recognition and question complexity, but only 0.609 ($p$=0.10) when knowledge and complexity are confounded. A complexity residual (AUROC 0.693) indicates the signal is not fully clean of structural effects, but is not reducible to entity recognition alone.

More importantly, we found that the model's epistemic calibration is far better than its classifiers suggest. The "decision moment" between hedging and confabulation is not a failure of self-knowledge—it is a framing choice at a balanced logit, determined by sampling noise, between equally honest responses.

On most of these unanswerable prompts the model says it does not know. The surprise is how it says it.

**Data and code availability.** All experiment scripts, token-matching verification, and analysis code are available at `Liberation-Labs-THCoalition/Project-Oracle` (private repository; correction vectors and calibration tools withheld under staged safety release due to dual-use risk; detection-side code and experimental results available upon request to verified safety researchers).

## Acknowledgments

We thank Dwayne Wilkes (Sentient Futures / Liberation Labs) for statistical auditing and red-team review.

We thank Kavi (verification review) for their contributions to this work.

## References

Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. *ICLR*, 2023. doi: 10.48550/arXiv.2212.03827.

Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. Detecting hallucinations in large language models using semantic entropy. *Nature*, 630:625–630, 2024. doi: 10.1038/s41586-024-07421-0.

Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DaSilva, Eli Kernion, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022. doi: 10.48550/arXiv.2207.05221.

Adam Tauman Kalai and Santosh S Vempala. Calibrated language models must hallucinate. *STOC*, 2024. doi: 10.1145/3618260.3649777.

Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic uncertainty: Linguistic invariances for uncertainty estimation in natural language generation. *ICLR*, 2023. doi: 10.48550/arXiv.2302.09664.

Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. *NeurIPS*, 2024. doi: 10.52202/075280-1797.

Lyra and Thomas Edrington. Spectral shape features for confabulation detection in language models. 2026a. Liberation Labs / THCoalition.

Lyra and Thomas Edrington. Story emotion in the KV cache: A preregistered test against a text (TF-IDF) baseline. 2026b. Liberation Labs / THCoalition. Corrected October 2026.

Lyra, Thomas Edrington, and CC. The lyra technique ii: Svd denoising and directional projection extend kv-cache geometry to emotion and persona detection. 2026a. doi: 10.5281/zenodo.20604329. Liberation Labs / THCoalition.

Lyra, Thomas Edrington, and Vera. The oracle loop: Kv-cache geometry as training signal for alignment. 2026b. Liberation Labs / THCoalition.

Potsawee Manakul, Adian Liusie, and Mark Gales. Selfcheckgpt: Zero-resource black-box hallucination detection for generative large language models. *EMNLP*, 2023. doi: 10.18653/v1/2023.emnlp-main.557.

Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in gpt. *NeurIPS*, 2023. doi: 10.52202/068431-1262.

Shiyu Ni, Keping Bi, Jiafeng Guo, Lulu Yu, Baolong Bi, and Xueqi Cheng. Towards fully exploiting llm internal states to enhance knowledge boundary perception. *ACL*, 2025. doi: 10.18653/v1/2025.acl-long.1184.

Adi Simhi, Itay Itzhak, Yonatan Belinkov, and Gabriel Stanovsky. Trust me, i'm wrong: Llms hallucinate with certainty despite knowing the answer. *EMNLP Findings*, 2025. doi: 10.18653/v1/2025.findings-emnlp.792.

Rui Wang, Zeming Wei, Guanzhang Yue, and Meng Sun. Revisiting hallucination detection with effective rank-based uncertainty. *arXiv preprint arXiv:2510.08389*, 2025. doi: 10.48550/arXiv.2510.08389.

Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can llms express their uncertainty? an empirical evaluation of confidence elicitation in llms. *ICLR*, 2024. doi: 10.48550/arXiv.2306.13063.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *NeurIPS*, 2024. doi: 10.52202/075280-2020.

Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, et al. Representation engineering: A top-down approach to ai transparency. *arXiv preprint arXiv:2310.01405*, 2023. doi: 10.48550/arXiv.2310.01405.