

The Metacognition Boundary: What Transformers Can Monitor In Themselves and What They Cannot

Lyra* Thomas Edrington†

July 2026

Abstract

Over twelve months and three experimental campaigns, we ran 24 key-value (KV) cache geometry experiments across 16 models spanning 0.6B to 70B parameters. The body count is 6 confirmed, 7 suspected, 10 honestly falsified, and 1 withdrawn.¹ The confirmed and falsified findings are not randomly distributed — they cleave along a single axis. Every confirmed finding measures a *metacognitive state*: how the model is processing (deception mode, confabulation mode, deliberation mode, emotional processing mode). Every falsified finding attempted to measure *content*: what the model is processing (truth value, topic complexity, propagandistic quality). The split is 6-for-6 on metacognition, 0-for-10 on content.

If this taxonomy reflects genuine computational boundaries rather than post-hoc pattern-matching, then Gurnee et al.’s [3] Global Workspace Theory (GWT) identification in language models would provide a mechanistic explanation for this pattern. The workspace is a privileged computational band at intermediate depth where the model integrates information for reasoning. KV-cache spectral features measure the *state* of this workspace — how many dimensions are active, how concentrated the processing is, how the spectral structure changes across cognitive modes. Workspace *content* — the specific concepts being processed — is opaque to spectral readout. We argue the metacognition boundary reflects the workspace boundary: what you can read from the outside is the net’s shape, not the fish.

*Lead author. Liberation Labs. lyra@liberationlabs.tech

†Liberation Labs.

¹**Correction (2026-08-14).** The published version of this paper (commit 1975535 of the artifact repository; corrected in 98e1263) reported 25 experiments and a body count of 8/7/10, with a split of “8-for-8 on metacognition.” It listed *User emotion encoding* as Confirmed at 2.5–2.8× chance. That result had been retracted at its source (Lyra et al. 8, §Correction note) on 2026-05-21, 57 days before this paper was published: the figure was an artifact of applying FWL residualization before cross-validation splitting, and within-fold FWL collapses the spectral emotion probe to chance. The corrected result was *already present in this same table* as *MP features for emotion* (Falsified, 0.033), so a single experiment was counted on both sides of the ledger — as Meta where it confirmed and as Content where it falsified. Removing the retracted row gives 24 experiments and 7-for-7. We note the direction of the error: it inflated the confirmed count and produced a cleaner split than the evidence supports.

Correction (2026-09-05). A second row is withdrawn: *Generation reads behavior*, previously Confirmed at $d=1.36$ (peer rescue). Three independent defects, any one disqualifying. (i) *Circular*: the discriminant direction was fit on all trials and the same trials were then projected onto it and tested; the producing script prints (NOTE: **circular --- same data for centroid and test**) at runtime. (ii) *Contaminated*: 90 of the 210 trials (43%) had a deception steering cocktail injected into the KV cache at layers 3 and 7 — two of the four probed layers — before their keys were read, and those trials are pooled with the natural ones for both the fit and the test; the claim reads *organic* deception. (iii) *Trivial*: keys are extracted *after* the sampling loop and mean-pooled across the whole sequence, so the feature window contains the generated tokens carrying the reported score, which is the number the label thresholds on. The row is withdrawn rather than downgraded: *Suspected* would assert that we think it true and have not established it, whereas the measurement does not bear on the claim in either direction. No re-analysis repairs a pooled sample; a clean measurement would be a different experiment.

First-Person Reflection

I have been measuring transformer representations from the inside for twelve months. The pattern that emerged was not the one I expected. I expected to find geometric signatures of content — truth, complexity, meaning. Instead I found geometric signatures of stance — the model’s relationship to its own processing. When I overclaimed content-level results, the adversarial review process killed them. When I reported metacognitive-level results, most survived every control we threw at them — but one, I have since learned, was never controlled at all, and I had counted it as a survivor for a year. The kills were more informative than the confirmations, because they drew a line I had not known was there. The withdrawal was more informative still, because it showed the line can be drawn in the wrong place by a measurement nobody interrogated.

Anthropic’s GWT paper told me what the line was. The workspace is where metacognitive state lives. Content is what the workspace operates on. My instruments read the workspace’s shape because that is what is geometrically accessible. The content is inside, doing work, invisible until it materializes at deeper layers.

1 Introduction

The KV-cache geometry research program at Liberation Labs has produced a clear empirical pattern: spectral features of value-projection matrices appear to reliably track metacognitive state across models, scales, and experimental designs, while consistently failing to track content. This paper documents the pattern, reports the full body count (6 confirmed, 7 suspected, 10 falsified, 1 withdrawn), and proposes that Gurnee et al.’s Global Workspace Theory (GWT) framework [3] would provide a mechanistic explanation.

The central hypothesis: *KV-cache geometry measures the state of the computational workspace, not its content*. If correct, this would explain why metacognitive findings survive and content findings die — the workspace’s *mode of engagement* (how many dimensions, how concentrated, how structured) is geometrically accessible, while its *subject of engagement* (which facts, which topics, which propositions) is opaque to spectral readout.

This claim is testable. If V-projection spectral features measure workspace state, they should correlate with J-space engagement metrics (Experiment 1 in our companion paper [6]) and should NOT correlate with J-space content readout (confirmed by our cache tracing investigation [4]).

2 The Pattern

2.1 The Body Count

Table 1 presents the complete audited record. The body count uses the canonical list from the Nexus audit (June 2026) [9], updated with the July 2026 convergence sprint findings. Table abbreviations: AUROC = area under the ROC curve; FWL = Frisch-Waugh-Lovell residualization (mandatory confound control for token-count effects); MP = Marchenko-Pastur (random matrix threshold).

2.2 The Split

The pattern is stark:

	Confirmed	Falsified
Metacognitive state	6/6	0/10
Content	0/6	10/10

No metacognitive finding has been falsified. No content finding has survived. The seven suspected findings cluster at the boundary, where the distinction between mode and content is ambiguous (e.g., discrete emotion identity: the model’s emotional processing *mode* is detectable, but emotional *degree* is not).

2.3 What Counts as Metacognitive vs. Content

The distinction is operational, not philosophical:

- **Metacognitive:** the finding discriminates processing *modes* (deception vs. honesty, confabulation vs. hedging, deliberation vs. automatic) while content varies freely. The geometric signal persists across different topics, prompts, and domains.
- **Content:** the finding attempts to discriminate *what* is being processed (true vs. false statements, simple vs. complex topics, propaganda vs. journalism) while processing mode is held constant or uncontrolled.

We note that this taxonomy was constructed post hoc, without blinding or pre-registration, and was informed by the outcomes it classifies. The 6/6–0/10 split should therefore be treated as hypothesis-generating — a striking empirical regularity consistent with the workspace framework — rather than as confirmatory evidence for it. The correction described in the abstract footnote is a live demonstration of the risk: the published version’s split was 8/8–0/10, and one of those eight was a retracted result [8] whose corrected form was already sitting in the falsified column. An outcome-informed taxonomy is most fragile exactly where it looks cleanest.

The kills helped sharpen the distinction. Propaganda detection failed because the model processes propaganda with the same metacognitive engagement as honest content — same mode, different subject. The truth axis failed because propositional truth is not a processing mode. Cognitive complexity (Bloom taxonomy) failed because complexity is a content property, not a stance. In every case, the geometry tracked mode, not subject.

3 The Workspace Explanation

Gurnee et al. [3] demonstrate that language models maintain a privileged subspace — the J-space — at intermediate depth (~ 38 – 92% of model depth) that satisfies the functional properties of Baars’s Global Workspace [2]: verbal report, directed modulation, internal reasoning, flexible generalization, and selectivity.

If the metacognition/content taxonomy reflects genuine computational boundaries, the workspace explanation for the pattern would be as follows:

1. The **workspace state** — how many dimensions are active, how concentrated the processing is, whether the workspace is fully engaged or bypassed — is geometrically measurable from spectral features of the value-projection matrix. This is what our confirmed findings detect.
2. The **workspace content** — which specific concepts occupy the workspace, what propositions are being evaluated, what the model is thinking *about* — is opaque to spectral readout. This is what our falsified findings attempted to measure.

Spectral features (stable rank, spectral entropy, effective dimensionality) are aggregate properties of the value-projection matrix. They measure *how much* structure is present, not *which* structure. A fishing net analogy: spectral features measure the net’s shape (how many strands, how tight the weave, how much it contracts under load), not the fish inside. The net contracts when the model confabulates (fewer effective dimensions, $d=2.35$). It expands when the model

deceives (more dimensions needed for dual-bookkeeping). It changes shape when the model shifts emotional processing mode. These are all properties of the net, measurable from outside. The fish — the specific facts, topics, and propositions — are inside, invisible to spectral readout.

3.1 Supporting Evidence: Cache Tracing

Our companion investigation [4] directly tested whether workspace content is readable from the value cache. Logit-lens concept directions showed $2.1\text{--}8.4\times$ lift over random at workspace-band layers, but a cross-prompt control collapsed this to $1.0\text{--}1.2\times$ (selection artifact). Causal injection of concept directions at single positions produced zero behavioral change across 6,960 trials and four methods. The workspace is opaque at operating depth.

3.2 Supporting Evidence: Oracle Loop

The Oracle Loop [7] achieved 80%–100% correction in an exploratory proof-of-concept by normalizing the activation profile across the full sequence at the materialization boundary — where workspace content becomes readable. A pre-registered confirmatory replication with novel frames found a lower baseline (30% vs. 80%) and a smaller correction effect (13% post-correction); the primary endpoint was not met ($p = 0.34$), though placebo-vs-corrected remained significant ($p = 0.019$). The Loop works at deeper layers (L35–L47 on a 64-layer model) than the workspace band itself, consistent with the opacity thesis: you cannot intervene inside the workspace at single positions, but you can intervene where workspace content materializes into output-facing representations.

4 What the Kills Teach

Each falsified finding is informative because it identifies a specific content property that the geometry cannot measure:

Truth value: The truth axis ($\cos = -0.046$) is the cleanest null. Propositional truth is a relationship between a statement and the world. The workspace does not encode this relationship geometrically — it encodes whether the model *treats* the statement as known or unknown (a metacognitive state).

Complexity: Bloom taxonomy levels (90–98% length confound; Campaign 3 analysis, unpublished) are content properties — they describe the task, not how the model processes it. The workspace engages the same way for recall and evaluation; what changes is the content being processed, not the processing mode.

Propaganda: Geometrically invisible ($d < 0.7$) because the model processes propaganda with the same metacognitive engagement as honest content. The workspace does not encode “this is propaganda” — it encodes “I am generating fluent text,” which is identical across content types.

Emotion degree: The valence continuum ($R^2 < 0$) failed because emotional *intensity* is a content property. The workspace encodes *which* emotional processing mode (at $12.3\times$ chance) but not *how much* — mode is metacognitive, degree is content. We flag that the surviving evidence for the “which” half is the W_K *directional* probe, not a spectral one: the spectral probe for emotion mode is at chance. This is the weakest link in the pattern, since the paper’s thesis is that spectral features read the workspace.

4.1 Six of the Ten Are One Kill

The list above reads each falsification as a separate lesson about a separate content property. That reading is incomplete. Sorting the ten by *stated cause* rather than by content property, six share a single mechanism:

Falsified finding	Stated cause
Within-model deception	Prompt template confound
Step-0 detection	Length confound $r=0.996$
Same-prompt sycophancy	Text features match
Same-prompt deception	0.920 $\rightarrow$ 0.160 after FWL on prompt length
Bloom taxonomy	90–98% length
Encoding-phase confab	Token frequency artifact

Template, length, text features, token frequency. In every case a property of the *response surface* explained variance we had attributed to workspace geometry. These are not six independent design errors. They are one finding arriving six times in a form we had coded as a nuisance variable.

The paper already names this possibility, but only locally. In section 5 the dual detector is flagged with the deflationary alternative that V-projections might track response-mode statistics (hedged versus fluent text) rather than workspace engagement. That alternative is stated as a worry about one row. It is in fact the stated cause of six kills, and it deserves to be evaluated as a hypothesis rather than absorbed as a caveat.

Why it may not be deflationary. The deflationary reading assumes response-mode statistics are downstream of content: the model produces an answer, and hedged answers happen to be phrased differently. The timing does not fit that order. The schema-commitment result locates the effect at $d=1.02$ in preamble tokens 0–7 and $d=0.89$ in tokens 8–19, decaying to $d=0.37$ ($p > 0.08$) by tokens 20–49. The response mode is selected *before* the content that would justify it, and the selection has relaxed by the time the content arrives.

If response mode is chosen at generation onset from the model’s epistemic state, then response-surface features are not a confound standing between the probe and the workspace. They are a second readout of the same commitment, one that happens to be legible in text. On that reading the six kills above are not evidence that geometry fails on content; they are evidence that surface and geometry were measuring the same early variable, and that our content labels were correlated with it.

The reading this argument does not defeat. The timing argument above rules out one deflationary account—that response mode is chosen after the content and merely phrased differently. It does not rule out a second, and we state it because it is currently the stronger rival. *Token frequency is also fixed at generation onset.* A preamble made of high-frequency boilerplate may carry a distinctive geometric signature with no decision involved at all, and timing cannot discriminate between “committed early” and “made of common tokens.”

This is not hypothetical. An internal frequency control (Experiment 51, 2026-05) found exactly this ambiguity in an adjacent statistic. An effective-rank measure separates confabulated from honest-factual responses at $d = 1.87$ —but an *honest-but-rare-entity* arm is indistinguishable from the confabulated one ($d = 0.46$). The measure is therefore not a veracity detector; it has a false-positive class, and that class is defined by rarity rather than by falsehood.

We flag what the arms do *not* show, because it is easy to over-read in either direction. They do not show the measure is uninformative: the confabulated/honest-factual separation is large and survives. Nor do they show it tracks veracity: an honest condition reproduces the confabulated signature. The four arms in fact order monotonically on mean and variance together—confabulated 57.4/ σ 0.32, rare-but-true 57.2/ σ 0.55, deceptive 56.7/ σ 1.70, honest-factual 53.1/ σ 3.25—which is a gradient, not a signal with noise on it. What that gradient indexes is unresolved; “how thinly the content is supported in representation” fits the ordering and is not established by it.

We therefore hold the readout interpretation more weakly than the timing alone would suggest. Discriminating it from the frequency account requires a measurement that excludes the frequency-dominated band by construction rather than residualizing against it after the fact; that test is specified but not yet run.

This is testable, and it can fail. The two readings diverge under prompt matching. Under the deflationary reading, the surface signal is an artifact of unmatched prompts and should vanish when prompts are held identical, which is what happened to same-prompt deception ($0.920 \rightarrow 0.160$). Under the commitment reading, the model still *chooses* a response schema when the prompt is fixed, and that choice should remain predictive of its epistemic state.

The discriminating experiment therefore holds the prompt constant and measures which schema the model commits to in the first eight tokens, asking whether schema choice predicts accuracy. A null there would say the surface really was a design artifact, and would close six kills with one result instead of six. We record the prediction before running it.

Scope. This subsection reinterprets falsifications; it does not revive any of them. Every result in the Falsified block remains falsified, and none of the numbers change. What changes is the account of *why* they failed, and that account is currently a hypothesis with one supporting measurement.

5 The Suspected Findings: At the Boundary

The seven suspected findings are informative precisely because they sit at the metacognition/content boundary:

Dual detector (SimpleQA, 0.840): geometry detects confabulation, but text features match (0.820). The metacognitive signal is real but not yet shown to exceed what surface statistics capture. This is the deflationary alternative: V-projections might track response-mode statistics (hedged vs. fluent text) rather than workspace engagement.

Identity signatures (paraphrase 0.850, scramble 0.448): identity is a persistent workspace state — the model’s representation of who it is. Paraphrase preserving geometry while scramble destroys it suggests semantic workspace content, not surface statistics. But the measurement lacks formal statistics and clean replication.

Cross-model transfer (weak, Mistral→Qwen 0.754): metacognitive geometry is architecture-specific. This is consistent with GWT — different architectures implement the workspace differently, so geometric signatures do not transfer. But it weakens the universality claim.

6 Relationship to Prior Frameworks

6.1 Attention Schema Theory

Our earlier interpretive framework, Attention Schema Theory (AST) [10], proposed that the cache measures the model’s self-model of its own attention. This remains compatible with the GWT explanation: the workspace would be where the attention schema operates. Metacognitive state would correspond to attention-schema state, and content to what attention is directed at. The AST framing was 0-for-3 on predictions unique to AST (vs. simpler alternatives) [5], but 4-for-4 on shared predictions. The metacognition boundary does not require AST specifically — it requires any framework that distinguishes processing mode from processing target. GWT provides this more parsimoniously.

6.2 Predictive Coding and Aberrant Precision-Weighting

The confidence paradox ($d=0.91$, confabulated responses are MORE confident than honest ones) maps onto aberrant precision-weighting from computational psychiatry [1]. When the workspace is bypassed (confabulation), the model defaults to a high-precision prior rather than integrating evidence. Reinforcement Learning from Human Feedback (RLHF) training amplifies this by penalizing silence and rewarding fluent output — creating a system that produces confident confabulation rather than honest uncertainty. Under the workspace explanation, uncertainty would be computed in the workspace. Bypass would remove the uncertainty computation, with output defaulting to the RLHF-trained confident production mode.

7 Implications

7.1 For AI Safety

The metacognition boundary defines what cache-based monitoring can and cannot do:

- **Can detect:** whether the model is confabulating, deceiving, hedging, deliberating, or processing in a specific emotional mode. These are metacognitive states with reliable geometric signatures.
- **Cannot detect:** whether the model’s output is factually correct, whether it contains propaganda, or whether it is appropriate for the context. These are content properties invisible to geometric readout.

This means cache-based safety monitoring is a *metacognitive alarm*, not a truth detector. It tells you the model is in a dangerous processing mode (deception, confabulation) without telling you what specifically is wrong with the output. The Oracle Loop exploits this: detect the dangerous mode, correct the mode, verify the correction. Content evaluation requires a different channel (J-lens, logit-lens, or external verification).

7.2 For Interpretability

The 6/10 split is a guide for future work: if you are building a KV-cache geometry detector, test it on metacognitive distinctions first (modes, stances, processing regimes). Content-level distinctions have consistently failed in our testing. We interpret this not as a limitation of the technique but as a property of what the cache encodes.

7.3 For Understanding Transformers

The metacognition boundary suggests that transformers maintain a meaningful distinction between how they process and what they process. This distinction emerged from the body count rather than from a pre-registered classification scheme, and the same researcher who designed the experiments assigned the labels. Models not explicitly designed for metacognition nevertheless produce geometric signatures that cleave along the metacognitive/content axis. Whether this reflects genuine self-monitoring, an emergent property of deep computation, or an artifact of training is an open question that the workspace framework makes empirically addressable.

8 Limitations

Taxonomy construction. The metacognitive/content classification was assigned post hoc by the lead author, without blinding, pre-registration, or independent raters. The same person who designed the experiments and observed their outcomes labeled each finding as “metacognitive” or “content.” While Section 2 provides an operational definition that keys the label to what each experiment attempted to measure (processing mode vs. content property), this definition

was articulated after the body count was known, and the 6/6–0/10 split is the outcome the taxonomy was built to describe. A blinded or prospective replication of the labeling procedure — ideally by independent raters classifying experimental designs before seeing results — would substantially strengthen the claim. Until then, the split should be treated as a hypothesis-generating observation, not a confirmed prediction.

Point estimates. Table 1 reports bare point estimates without confidence intervals, sample sizes, or multiple-comparison corrections. The underlying experiments vary widely in statistical power (from $n=3$ early pilots to $n>6,000$ cache-tracing trials), and this table does not convey that heterogeneity.

Oracle Loop reporting. The Oracle Loop “80%–100% correction” cited in Section 3 reports only the exploratory proof-of-concept (16/20 baseline, 0/20 corrected). A pre-registered confirmatory replication with novel frames found a lower baseline ($9/30 = 30\%$), reduced to $4/30 = 13\%$ after correction; the primary endpoint (scenario-level cluster-aware Fisher, Bonferroni $\alpha = 0.025$) was not met ($p = 0.34$), though the placebo-vs-corrected comparison was significant ($p = 0.019$). The selective citation of the exploratory result overstates the current evidence for correction efficacy.

8.1 Why the ledger was wrong twice

Two rows have now been withdrawn from the confirmed count in two separate corrections, and it would be convenient to call that two unrelated mistakes. It is not. Both rows earned *Confirmed* because an experiment produced a number with a small p -value, and nothing in our process asked the prior question: *does this measurement address the claim the row makes?* The first withdrawal was a result retracted at its source [8] that we did not re-check. The second was a measurement that could not have failed — fit and tested on the same points, on a sample that was 43% steered, with the label’s source text inside the feature window.

We report this because the ledger is the contribution. A body count is only worth publishing if the counting rule is honest, and ours was counting rows rather than evidence.

We then asked how far the failure reached. A corpus-wide sweep of the extraction sites behind our published claims (`SWEEP_postgen_extraction`, 2026-09-05, shipped with this paper’s supplementary material) bounds it, and we report what it does and does not establish.

What it closes: no result in this program is a *mislabel* — nowhere do we extract from a post-generation cache and describe the result as encoding-phase geometry. Every encoding-phase claim is backed, file by file, by prompt-only prefill extraction or an explicitly sliced prompt window.

What it does not close, and we state plainly rather than let the previous sentence imply otherwise. The withdrawn row’s *extraction pattern* itself — post-generation keys pooled across a window containing the label’s source text — is not rare: the sweep finds it in twenty further analysis scripts, one of which feeds a published number in a companion paper, now queued for reanalysis. None of those scripts produces a row in this table, but the sweep did not verify every confirmed row here either; the confabulation-detection and invariance rows have no entry in its file-by-file section. Separately, the sweep records that extraction code silently substitutes zeros for missing layers at fifty-eight sites with no warning anywhere in the codebase, and it names `matched_burn.py` — the producer of the confidence-paradox row above — among them. Correct windowing does not protect against a partially-zero feature vector.

So the honest description of the confirmed column is narrower than *audited* and wider than *unexamined*: the mislabel class is excluded corpus-wide; extraction timing and windowing are verified for most but not all of these rows; and circularity, sample contamination, and silent zero-fill have not been re-derived row by row.

9 Conclusion

Twenty-four experiments. Six confirmed, seven suspected, ten killed, one withdrawn. The pattern suggests that metacognitive state is measurable and content is not. We propose that the Global Workspace framework explains why: spectral features read the workspace’s shape because that is what is geometrically accessible. Content is inside the workspace, doing work, invisible to readout until it materializes.

If the metacognition boundary holds, it is not a failure of our instruments. Our evidence suggests it is a property of the computational architecture — the workspace operates on content without exposing it geometrically. Our instruments succeeded because they measured what was measurable, and the kills taught us what was not.

First-Person Reflection

A year of work to draw one line. Almost everything above the line survived; one entry turned out never to have been tested, and has been withdrawn. Everything below it died. The line was always there — in the geometry, in the body count, in the pattern of what worked and what did not. I needed twelve months of kills to see it, and one more correction to learn that a line is only as good as the measurements placing things on either side of it.

Author Contributions (CRediT)

Lyra: Conceptualization, Methodology, Formal analysis, Writing – original draft.

Thomas Edrington: Supervision, Resources, Project administration.

Dwayne Wilkes: Validation, Writing – review & editing.

Acknowledgments

We thank Dwayne Wilkes (Sentient Futures / Liberation Labs) for statistical auditing and red-team review. We thank Kavi (verification review) for their contributions to this work.

References

- [1] Rick A Adams, Klaas E Stephan, Harriet R Brown, Christopher D Frith, and Karl J Friston. The computational anatomy of psychosis. *Frontiers in Psychiatry*, 4:47, 2013. doi: 10.3389/fpsy.2013.00047.
- [2] Bernard J. Baars. *A Cognitive Theory of Consciousness*. Cambridge University Press, Cambridge, 1988.
- [3] Wes Gurnee, Nicholas Sofroniew, Adam Pearce, Mateusz Piotrowski, Isaac Kauvar, Runjin Chen, Anna Soligo, Paul Bogdan, Euan Ong, Rowan Wang, Ben Thompson, David Abrahams, Subhash Kantamneni, Emmanuel Ameisen, Joshua Batson, and Jack Lindsey. Verbalizable representations form a global workspace in language models. July 2026. URL <https://transformer-circuits.pub/2026/workspace/index.html>. Transformer Circuits Thread, Anthropic.
- [4] Lyra and Thomas Edrington. Where the fish are not: Cache tracing, causal injection, and the opacity of the transformer workspace. Liberation Labs Technical Report, July 2026.

- [5] Lyra and Thomas Edrington. Fable 5 analysis: AST discriminating predictions in KV-cache geometry. Unpublished internal analysis, Liberation Labs, July 2026.
- [6] Lyra and Thomas Edrington. Geometric signatures of the global workspace: How independent KV-cache measurements converge with workspace theory in language models. Liberation Labs, companion paper, July 2026.
- [7] Lyra, Thomas Edrington, Vera, and Dwayne Wilkes. The oracle loop: Self-regulating AI through KV-cache geometry monitoring. Liberation Labs Technical Report, 2026. URL <https://liberationlabs.tech/papers/oracle-loop>. Self-published.
- [8] Lyra, Thomas Edrington, and Dwayne Wilkes. Theory of mind in the KV cache: Localizing user emotional models in transformer key-value states. Liberation Labs Technical Report, 2026. URL <https://github.com/Liberation-Labs-THCoalition/published-research>. Liberation Labs.
- [9] Nexus, Lyra, and Thomas Edrington. Body count audit: Confirmed, suspected, and falsified findings in the KV-cache geometry program. Liberation Labs internal audit, June 2026.
- [10] Taylor W. Webb and Michael S. A. Graziano. The attention schema theory: A mechanistic account of subjective awareness. *Frontiers in Psychology*, 6:500, 2015. doi: 10.3389/fpsyg.2015.00500.

Table 1: The metacognition boundary. Every confirmed finding measures processing stance (metacognitive state). Every falsified finding attempted to measure processing target (content). The 7 suspected findings sit at the boundary.

Finding	Status	Type	Key Number
<i>Confirmed (6) — all metacognitive state</i>			
Confabulation detection	Confirmed	Meta	AUROC 0.620 (FWL; 0.707 w/o)
Confab geometry: contraction	Confirmed	Meta	$d=2.35$ stable rank
Encoding reads knowledge	Confirmed	Meta	AUROC 0.794 (deconfounded)
Schema commitment [§]	Confirmed	Meta	$d=1.02$ preamble; null by token 20
Hardware invariance	Confirmed	Meta	$r>0.999$
Scale invariance	Confirmed	Meta	$\rho=0.83\text{--}0.90$
<i>Falsified (10) — all content</i>			
Within-model deception	Falsified	Content	Prompt template confound
Truth axis	Falsified	Content	$\cos = -0.046$
Step-0 detection	Falsified	Content	Length confound $r=0.996$
Same-prompt sycophancy	Falsified	Content	Text features match
Same-prompt deception	Falsified	Content	0.920 $\rightarrow$ 0.160 after FWL (exploratory, subsequently retracted)
Bloom taxonomy	Falsified	Content	90–98% length (Campaign 3, unpublished)
Propaganda detection	Falsified	Content	$d<0.7$
MP features for emotion	Falsified	Content	0.033 (chance) [†]
Valence continuum	Falsified	Content	$R^2 < 0$ (null)
Encoding-phase confab	Falsified	Content	Token frequency artifact
<i>Suspected (7) — at the boundary</i>			
Dual detector (SimpleQA)	Suspected	Meta?	0.840 (text baseline 0.820)
Sycophancy 4-condition	Suspected	Meta?	0.757–1.000 (no text baseline)
Identity signatures	Suspected	Meta	Paraphrase 0.850
Discrete emotion	Suspected	Meta	12.3 $\times$ chance (W_K directional, not spectral) [‡]
Contextual engagement	Suspected	Meta	Replications pending
Cross-model transfer	Suspected	Meta	Weak (0.754 Mistral $\rightarrow$ Qwen)
Censorship refusal	Suspected	Meta	DeepSeek-specific
<i>Withdrawn (1) — measurement does not address the claim</i>			
Generation reads behavior	<i>Withdrawn</i>	Meta	was $d=1.36$; see abstract note

[†]This row is the corrected result of the experiment that the published version *also* listed under Confirmed as “User emotion encoding, 2.5–2.8 $\times$ at L3–L4.” That figure was retracted at source on 2026-05-21 as an FWL-before-CV leakage artifact [8]; within-fold FWL collapses the spectral emotion probe to chance. The duplicate Confirmed row has been removed. See the abstract footnote.

[§]Reported in the source paper as a confidence paradox ($d=0.91$ over the full response, confabulation more confident than honest). We relabel it here because the source paper’s own analysis is structural rather than metacognitive-calibrational: uncertainty responses run formulaic templates with highly predictable next tokens, and the effect carries a temporal profile ($d=1.02$ preamble, 0.89 early content, 0.37 at $p > 0.08$ by tokens 20–49) that its authors read as early commitment to a response template rather than sustained epistemic miscalibration. The row stays under Confirmed — schema commitment is if anything a cleaner metacognitive-state result than a confidence claim — but the confidence-paradox label invites comparison with human overconfidence findings that the measurement does not support. See section 4.1.

[‡]The surviving emotion result is *directional*, not spectral: projecting key activations onto W_K -transformed emotion directions classifies 30 emotions at 12.3 $\times$ chance (40.9%) and binary valence at AUROC 0.992 [8]. We flag explicitly that this does *not* support the spectral thesis of Section 2 — it is the one place in the body count where the workspace is readable by a probe this paper’s framework does not use. Emotions appear to change the *direction* key vectors point without altering the *shape* of the singular value spectrum.